

## Highlights

### **Tangent-space features for longitudinal EEG movement-intent decoding, and the protocol condition that bounds them**

Esm E Moula Chowdhury Abha, Md Wahiduzzaman Suva, Muhammad Hasibur Rashid Chayon

- A tangent-space branch supplies the channel covariance a log-power stem discards.
- Accuracy rises on four of five cohorts, 51 participants, three seeds each.
- It beats eight published decoders trained in the same pipeline on each cohort.
- Its label-free reference fails where a class block outlasts the estimator memory.
- Block length against estimator memory, read from the protocol, separates all five.

# Tangent-space features for longitudinal EEG movement-intent decoding, and the protocol condition that bounds them

Esm E Moula Chowdhury Abha<sup>a</sup>, Md Wahiduzzaman Suva<sup>a,\*</sup> and Muhammad Hasibur Rashid Chayon<sup>a</sup>

<sup>a</sup>Department of Computer Science, American International University-Bangladesh, Dhaka, Bangladesh

## ARTICLE INFO

### Keywords:

brain-computer interface  
electroencephalography  
test-time adaptation  
covariate shift  
selective prediction  
calibration  
lower-limb exoskeleton

## ABSTRACT

Compact EEG decoders summarise each window as log power per channel, which is the diagonal of its spatial covariance. How channels covary is discarded, and on movement-intent tasks that structure carries information the diagonal does not.

We add a tangent-space branch beside the log-power stem. Each window's covariance is shrunk, mapped to the tangent space at a reference the network maintains from the input alone, and projected into the classifier. No labels are used at inference and the reference keeps updating on the test stream. Across five cohorts and 51 participants, three seeds each, the branch raises balanced accuracy on four: +0.064 on a 20-participant motor-execution cohort held out from all development ( $p = 8.8 \times 10^{-5}$ , higher in 20 of 20), +0.021 on a lower-limb exoskeleton cohort ( $p = 0.047$ ), +0.040 on BCI Competition IV-2a ( $p = 0.020$ ), and +0.045 on a second exoskeleton cohort. It beats all eight published decoders trained in the same pipeline on each cohort where the comparison was run.

On the fifth cohort it costs 0.175 and falls near chance, and that failure is predictable before training. The branch's reference is a running mean over the incoming windows, so it is safe only while a single covariance update spans more than one class block. A block lasts 214 windows there against an update of 32; on the four cohorts where the branch helps a block is 1 to 18 windows and the memory spans many blocks. Both quantities are fixed before training, the block length by the protocol and the memory by the optimiser settings, and the comparison separates all five cohorts. Moving the memory across the block length on a single cohort, with nothing else changed, switches the effect off and on: it is +0.045 at a memory of 160 windows, +0.048 at 40, and -0.076 at 8, worse in every participant. Enlarging the memory on the failing cohort recovers 74 % of the loss, which is causal support for the account without restoring parity. A stronger prescription, setting the adaptation rate from the block length, was pre-registered and falsified.

## 1. Introduction

Spinal cord injury leaves roughly 250,000 to 500,000 people per year with impaired or absent lower-limb function, and powered exoskeletons driven by non-invasive brain signals are one route to restoring overground mobility [1–3]. Decoding the wearer's intent to walk or to stop from scalp EEG is the control problem at the centre of such systems [4, 5]. It is not well described by per-epoch classification accuracy alone, for three reasons that shape the design presented here.

*Longitudinal drift* The decoder is fitted once and must keep working weeks later. In the primary cohort used here this is literal: the model is fitted on sessions 1–3 and evaluated on sessions 4–9, recorded across weeks. Electrode montage, impedance and skin contact all change between sessions, and the resulting shift in second-order statistics is the dominant source of non-stationarity in EEG decoding [6–8]. Session drift degrades fixed spatial filters and is the reason most BCI systems require a per-session calibration block [9].

*Asymmetric cost* A false Walk command moves a person's legs. Balanced accuracy scores that identically to a missed Walk and cannot distinguish one sustained false activation from fifty scattered false windows. The quantity a wearer experiences is spurious activations per minute of standing, which we report alongside accuracy throughout.

*What a log-power stem cannot see* Compact EEG decoders reduce each window to log power per channel after spatial filtering [10]. That summary is the diagonal of the window's spatial covariance. The off-diagonal entries, how channels covary, are discarded, and covariance structure is the basis of the Riemannian methods that remain competitive on movement-intent tasks [11, 12]. Whether a convolutional decoder is losing usable information at that step is an empirical question, and it is the one this paper answers.

Existing responses to drift are mostly applied offline in preprocessing. Euclidean Alignment whitens each recording by the inverse square root of its mean covariance [13]; Riemannian Procrustes analysis and parallel transport generalise this to the manifold of symmetric positive-definite (SPD) matrices [14–16]. All of these estimate the alignment statistic over a complete recording, which means the target recording must exist before the decoder can run on it. Test-time adaptation removes that requirement in computer vision by updating normalisation statistics or parameters online [17–19], and recent work shows that such methods degrade

\*Corresponding author

✉ 26-94089-2@student.aiub.edu (E.E.M.C. Abha);

26-94088-2@student.aiub.edu (M.W. Suva); chayon@aiub.edu (M.H.R. Chayon)

ORCID(s):

when the test stream is temporally correlated instead of i.i.d. [20–22]. EEG data acquired under a blocked experimental protocol is temporally correlated by construction, which is the connection this paper makes precise.

Our contributions are:

1. Add a tangent-space branch beside the log-power stem: each window's covariance is shrunk, mapped to the tangent space at a reference estimated from the input alone, and projected into the classifier, with no labels used at inference (Section 4.2).
2. Measure it on five cohorts and 51 participants with three seeds each. Accuracy rises on four, including +0.064 in 20 of 20 participants of a cohort held out from every design decision (Section 6.9).
3. Compare against eight published decoders trained in one pipeline with shared splits, seeds and metric code, and report the cohort where most of them fail to train rather than letting the resulting margins stand (Section 6.10).
4. Identify the branch's failure mode and predict it from the protocol. The reference is a running mean over incoming windows, so it is safe only while its memory outlasts a class block; on the cohort where a block outlasts that memory the branch costs 0.175 and falls near chance (Condition 1, Section 6.6).
5. Test that account causally by enlarging the covariance update at fixed block length, which recovers 74 % of the loss without restoring parity (Section 6.9).
6. Pre-register a stronger version of the same condition, one that sets the adaptation rate from the block length, and report it as falsified (Section 6.7).

The branch costs parameters. It adds 238,028 to a 19,505-parameter stem path, plus 32,896 for the fusion it requires, so the reported model is no longer compact and we do not claim parameter efficiency; Section 6.11 reports the counts beside the accuracies so the trade is visible.

## 2. Related work

### 2.1. Compact convolutional decoders for EEG

ShallowFBCSPNet and Deep4Net established that a square  $\rightarrow$  pool  $\rightarrow$  log pipeline over learned spatial filters recovers the band-power features that filter-bank common spatial patterns computes explicitly [10]. EEGNet reduced the parameter count by an order of magnitude using depthwise-separable convolutions [23], and EEGNeX refined the same family [24]. These architectures work well within a session and are the standard baselines in the field [25, 26]. Their limitation for our setting is that the learned spatial filters are fixed at training time, so a montage or impedance change between sessions is not compensated at all. We respond by placing an adaptive whitening stage before the spatial filters, so that what the filters see is referred to a covariance estimated on the current recording (Section 4.3).

### 2.2. Attention-based decoders

EEG Conformer pairs a ShallowFBCSP-style stem with a transformer encoder [27], and EEG-TCNet applies temporal convolutions to the same problem [28]. The stem-plus-transformer combination is already published: an independent implementation of it in our pipeline reproduced EEG Conformer's score without improving on it (0.848 against 0.843), and we therefore claim no novelty for that pairing. What these models share with the convolutional family is a fixed front end, which is the property the present work targets.

### 2.3. Alignment and transfer in BCI

Riemannian methods treat the spatial covariance as a point on the SPD manifold and classify with geodesic distances [11, 12, 29]. For transfer, Euclidean Alignment whitens each subject or session by  $\mathbf{M}^{-1/2}$  where  $\mathbf{M}$  is its mean covariance [13]; Riemannian Procrustes analysis adds rotation and scaling [14], and parallel transport moves covariances between tangent spaces [15, 16]. These methods are effective and label-free, and reviews report them as the strongest general-purpose transfer tools for EEG [30, 31]. Their limitation is operational. The alignment statistic is computed over a whole recording, so the recording must be complete before decoding starts, and within-recording drift is not tracked. We respond by moving the estimator inside the network as a layer that updates online, and we then show what that move costs (Section 6.6).

### 2.4. Riemannian and tangent-space representations

Spatial covariances are symmetric positive-definite matrices and do not form a vector space, so the methods built on them operate on the SPD manifold. The standard recipe computes a covariance per window, maps it to the tangent space at a reference point, and classifies the resulting vector with a linear model [11, 12, 29]. The reference is usually the Riemannian or arithmetic mean of the training covariances, and the choice matters because the map is only a faithful local approximation near that point. These representations remain competitive with convolutional decoders on movement-intent tasks and are the strongest general-purpose transfer tools reported in reviews [30, 31]. SPD networks learn the manifold operations instead of fixing them [32].

Two differences separate the branch used here from that line. The tangent map is a second input to a convolutional decoder rather than the representation the decoder is built on, so first-order band power and second-order structure are read together instead of one replacing the other; Section 6.9 reports that the two are complementary and that the branch alone is weaker than the stem alone. And the reference point is estimated from the incoming windows at inference rather than fixed after fitting, which is what makes the representation track a changing session and is also, as Section 6.6 shows, what makes it fail under a blocked protocol.

## 2.5. Test-time adaptation

Updating normalisation statistics on the test distribution improves accuracy under image corruption in vision [18, 19], and Tent extends this to updating affine parameters by entropy minimisation [17]. Test-time training adds a self-supervised objective at inference [33]. A recurring finding in this literature is that these methods assume the test stream is i.i.d., and degrade when it is not: NOTE addresses temporally correlated streams by maintaining a class-balanced reservoir [20], SAR stabilises adaptation under distribution shift with small or imbalanced batches [21], and CoTTA and LAME address error accumulation in continual adaptation [22, 34]. Our contribution relative to this line is quantitative: we show that in a blocked EEG protocol the relevant comparison is between the estimator's effective memory and the duration of a single-class block, and that this comparison predicted a new cohort in advance. The constraint should also have an upper side, because the same estimator is what tracks the drift the layer exists to remove; that side is not measured here (Section 8).

## 2.6. Calibration and selective prediction

Modern networks are overconfident and expected calibration error [35, 36] degrades further under distribution shift [37], which is the regime a longitudinally deployed decoder operates in. Selective classification allows a model to abstain [38, 39], and SelectiveNet trains the classifier and the selection function jointly under a coverage constraint [40, 41]. A device that must not move a wearer's legs on weak evidence is an obvious candidate for a reject option, so we built one in and tested it. It did not work: the trained selection function ranked windows no better than the classifier's own confidence, and the head is not part of the reported model. Section 6.2 gives the measurement. We claim no contribution to selective prediction and report calibration error throughout because it is what a controller would threshold on.

## 2.7. EEG correlates of gait and exoskeleton control

Walking modulates cortical activity in a gait-phase-locked manner [42, 43], and treadmill walking can be decoded from scalp EEG [5, 44]. Robot-assisted gait modulates midline sensorimotor rhythms [45]. The intent signal our decoder targets is event-related desynchronisation in the mu and beta bands, a power decrease over sensorimotor cortex preceding and accompanying movement [46, 47]. This motivates a representation in which log-power ratios are preserved (Section 4.1), because ERD is defined as a relative power change.

## 3. Problem formulation

Let  $\mathbf{X} \in \mathbb{R}^{C \times T}$  be a window of  $C$  channels and  $T$  samples with label  $y \in \{\text{Stop}, \text{Walk}\}$ . A recording session  $s$  supplies windows drawn from  $p_s(\mathbf{X}, y)$ . We assume the conditional  $p(y | \mathbf{X})$  is stable across sessions but the marginal  $p_s(\mathbf{X})$  is not, which is the covariate-shift assumption standard

in this literature [6, 48]. Concretely we take the shift to act primarily on the spatial covariance  $\Sigma_s = \mathbb{E}_{p_s}[\mathbf{X}\mathbf{X}^\top]/T$ .

The decoder is fitted on sessions  $S_{\text{fit}}$  and evaluated on disjoint later sessions  $S_{\text{test}}$ , with no labels available from  $S_{\text{test}}$  at any point. This is the constraint that rules out per-session recalibration and motivates a label-free correction.

We quantify session drift by the affine-invariant Riemannian distance between SPD matrices,

$$\delta(\mathbf{A}, \mathbf{B}) = \left\| \log(\mathbf{A}^{-1/2} \mathbf{B} \mathbf{A}^{-1/2}) \right\|_F, \quad (1)$$

in preference to the Frobenius distance, because covariances do not form a vector space and the Frobenius distance is dominated by overall scale, which would flatter any whitening operation [11, 12].

Two timescales govern the analysis in Section 6.6. Let  $\tau_{\text{blk}}$  be the median duration, in windows, of a contiguous single-class block in the recording protocol, and let  $\tau_{\text{drift}}$  be the timescale over which  $\Sigma_s$  changes appreciably. Both are properties of the experimental design and the recording hardware, independent of the model.

## 4. Method


The reported decoder has two components. A multi-scale log-power stem reads the diagonal of each window's spatial covariance, and a tangent-space branch reads the rest of it. Their outputs are concatenated and passed to one linear classifier. The stem path carries 19,505 parameters and the branch 238,028, with a further 32,896 in the fusion layer the branch requires, so the branch is the larger part of the model by an order of magnitude and Section 6.11 reports what that buys.

Three further components were built and tested, and none is retained: an adaptive alignment layer, a causal cross-epoch transformer and a selective head. Section 4.3 describes them and points at the measurements that removed them, because each defines an arm of the ablation and one of them, the alignment layer, supplies the mechanism that explains where the branch fails.

### 4.1. Multi-scale power stem

Three parallel temporal resolutions (64, 128 and 256 ms), each with its own depthwise spatial filters, are followed by square  $\rightarrow$  pool  $\rightarrow$  log, after Schirrmeister et al. [10]. Amplitude normalisation precedes the squaring so that log-power ratios survive: scaling a window by  $s$  shifts its log-power by  $2 \log s$ , a constant that a later affine layer absorbs, whereas normalising after the square destroys the ratio in which ERD is defined [46].

This ordering is load-bearing, and the evidence is a dose-response across representations that differ in how much marginal power they carry (Table 7). Removing per-window amplitude normalisation is worth +0.134 and +0.174 on band-power features, which are pure marginal power; -0.016 and +0.029 on tangent-space features, which carry power only through the covariance diagonal; and -0.004 and +0.000 on a correlation-only representation,



Reported model: 290,943 parameters. The branch is 238,028 of them, and the fusion it requires a further 32,896: 93 % together.

*Built and measured, not retained:*

Adaptive alignment  
whitens by  $M^{-1/2}$   
substitute for the branch

Cross-epoch transformer  
594,816 parameters  
costs accuracy

Selective head  
coverage-constrained  
no usable reject option

**Figure 1:** The reported decoder. A window enters both paths. The stem applies three temporal resolutions with depthwise spatial filters, squares, pools and takes the logarithm, giving log band power per channel. The branch forms the window's covariance, shrinks it, maps it to the tangent space at a running reference and projects the result. The reference update (dashed) uses no labels and stays active at inference, which is what lets the representation follow a changing session and also what makes it fail where a class block outlasts one update (Section 6.6). Parameter counts are exact and sum to the total.

which is scale-free by construction and therefore cannot respond. A representation that carries no marginal power shows no effect, which is what identifies the mechanism instead of merely demonstrating it.

These are participant-paired comparisons, so they can be tested as such. Wilcoxon signed-rank tests over participants, Holm-corrected across the three representations within each cohort, give band power +0.174 (better in 7 of 8,  $p_{\text{Holm}} = 0.047$ ) on cohort B and +0.134 (6 of 7,  $p_{\text{Holm}} = 0.094$ ) on cohort A; tangent space +0.029 (5 of 8,  $p_{\text{Holm}} = 0.500$ ) and -0.016 (2 of 7,  $p_{\text{Holm}} = 0.750$ ); and the correlation-only control +0.000 (3 of 8,  $p_{\text{Holm}} = 1.000$ ) and -0.004 (3 of 7,  $p_{\text{Holm}} = 0.750$ ). **The effect reaches significance on the larger cohort and is null in the scale-free control, as the mechanism predicts;** cohort A is consistent in sign but underpowered at  $n = 7$ , where the smallest attainable Holm-corrected  $p$  is 0.047.

A separate sweep over where normalisation is applied in the network confirms the ordering at the architecture level: on cohort A with 4 s windows, per-window normalisation scores 0.849, global per-recording normalisation 0.878, and

no normalisation 0.864, so the global setting used here is the best of the three and per-window normalisation is the worst.

Table 8 repeats the normalisation comparison across six published decoders, in an earlier training pipeline. Per-window normalisation is the best choice for two of them and the worst for two, and most differences are below the 0.02 noise threshold, so this comparison neither supports nor contradicts the representation-level dose-response in Table 7. Table 7 also reports the control that retired an amplitude side-channel: supplying a per-window amplitude feature and supplying the same feature with its window assignment shuffled differ by less than 0.02 for four of five models.

## 4.2. Tangent-space branch

The stem of Section 4.1 reduces each window to log power per channel, which is the diagonal of its spatial covariance. The off-diagonal structure, how channels covary, is discarded. The branch restores it. For a window  $X$  we form the per-window covariance  $C = XX^T / T$ , normalise its trace, shrink it toward a scaled identity by  $\lambda = 0.1$ , and

map it to the tangent space at a reference  $M$ ,

$$\phi(X) = \text{vec}(\log(M^{-1/2}CM^{-1/2})), \quad (2)$$

with off-diagonal entries scaled by  $\sqrt{2}$  so that the vector's Euclidean norm equals the matrix Frobenius norm. A linear projection maps  $\phi(X)$  to the embedding width and concatenates it with the stem before the classifier, which leaves the classifier and the selective head unchanged in shape so the branch ablates as one module.

The reference  $M$  is a running mean over the covariances of incoming windows, updated by (4) with momentum  $m$  and no labels, and it keeps updating at inference. This is what lets the representation follow a session as it changes, and it is also the component that fails: because  $M$  is an average over whatever is currently streaming, a protocol that presents one class for longer than a single update spans will drive  $M$  onto that class. Condition 1 states when that cannot happen and Section 6.9 measures it on five cohorts. The condition was derived on the alignment layer of Section 4.3, which shares this estimator, and it transfers to the branch unchanged.

The tangent map carries no gradient; only the projection is learned. Backward through an eigendecomposition contains  $1/(\lambda_i - \lambda_j)$  terms that diverge on near-degenerate spectra, and covariances estimated from 60 channels and 400 samples are routinely near-degenerate. Nothing measured is lost by detaching it: the probe that motivated the branch was itself a linear map on fixed tangent features, and that is what is trained here, jointly with the network rather than after it. The branch carries 238,028 parameters and the fusion layer it requires a further 32,896, against 19,505 in the stem path; the reported model is 290,943 in total, of which the branch and its fusion are 93 %.

### 4.3. Components tested and not retained

Three components were built, measured, and left out of the reported model. Each is an arm of the ablation in Section 6.2, and they are described here because a reader cannot judge an ablation against arms that are never defined.

*Adaptive alignment* `abelsec:align` The layer maintains a running mean spatial covariance  $M$ , applies  $M^{-1/2}$ , and interpolates between the aligned and the raw signal through a learned scalar gate  $g \in [0, 1]$ :

$$X' = g M^{-1/2}X + (1 - g)X. \quad (3)$$

Each window's covariance is normalised by its own trace before being accumulated, so that the estimate is an average of shapes instead of an average weighted by power:

$$M \leftarrow (1-m)M + \frac{m}{N} \sum_{n=1}^N \frac{C_n}{\text{tr}(C_n)/C}, \quad C_n = \frac{\tilde{X}_n \tilde{X}_n^T}{T}, \quad (4)$$

where  $\tilde{X}_n$  is the mean-centred window,  $N$  the batch size and  $m$  the momentum. Without the trace normalisation the estimate is power-weighted. On a cohort whose majority class carries 1.52× the power of the minority class, the unnormalised estimate sits at Riemannian distance 0.306 from

the majority-class covariance, that is, it is the majority class; the normalised estimate sits at 3.899 and 3.194 from the two class covariances, between them. Section 6.5 reports what correcting this did and did not achieve.

Two properties separate this layer from batch normalisation [49]. The update in (4) uses only the covariance of incoming windows and therefore runs on unlabelled test data. And the estimate continues to update in evaluation mode, whereas BatchNorm freezes its running statistics at test time. Freezing is correct when train and test share a distribution; here they are weeks apart, and freezing discards the adaptation the layer exists to provide.

*Propagation of the aligned signal* A whitening stage placed before a normalisation layer invites the objection that the normalisation simply undoes it. With identical initialisation on held-out windows, stem outputs computed with and without alignment differ by 36 % in relative magnitude and their correlation structure differs by 0.108, so the transform is not absorbed downstream.

This layer is not in the reported model. On the cohort where the branch works best it changes accuracy by less than the spread between seeds, and on the cohort used for development the branch and the layer are substitutes rather than complements, since both read the same second-order structure; Section 6.9 gives the measurement. Its estimator is retained, as the branch's reference, and the rate condition derived on it is what predicts where the branch fails.

*Causal cross-epoch context* A causal transformer encoder [50] over a sequence of  $K = 8$  epochs at stride 3 (a 14.5 s span), following the epoch-sequence formulation standard in sleep staging [51–53]. All masks are causal, so only present and past information is used, which is the only variant a worn device can run. Measured state persistence in cohort A supports a sequence model: the self-transition probability is 0.984 for Stop and 0.955 for Walk at a 0.5 s step, corresponding to dwell times of 12 to 32 s. This pathway is not in the reported model: it accounts for 594,816 parameters and costs accuracy on the development cohort.

*Selective head* A classifier paired with a selection function trained under a coverage constraint [40], with a full-coverage auxiliary term so that the backbone continues to learn from every sample, including the declined region. We use the formulation as published. This head is not in the reported model. The trained selection function did not rank windows better than the classifier's own confidence, so it provided no usable reject option (Section 6.2).

## 5. Experimental setup

### 5.1. Data

**Cohort A (exoskeleton, OpenNeuro ds007788).** Seven participants, nine sessions each, recorded across weeks while walking in a powered lower-limb exoskeleton [4]. Sixty-channel active EEG at 100 Hz; walk/stop labels are taken from the exoskeleton controller state. Windows of

4 s at 0.5 s stride give approximately 5031 fitting windows per participant. The model is fitted on sessions 1–3 and evaluated on sessions 4–9.

**Cohort B (treadmill MoBI).** Eight participants, three trials each, 64-channel EEG at 100 Hz, walk/stand from treadmill phase [44]. Fitted on trial 1 and evaluated on trials 2–3, with 2 s windows. The cohort is from an independent laboratory with different sensors, a different task framing, and reversed class balance: Stop is the minority class at approximately 12 % here and the majority class in cohort A.

**Cohort C (motor execution, PhysioNet EEGMIDB).** Twenty participants, S001–S020, from the PhysioNet motor movement and imagery database [54, 55], recorded with 64-channel EEG at 160 Hz and resampled to 100 Hz. Only the motor-execution runs are used, and the task is rest against movement of the hands or feet. The model is fitted on runs 3, 5, 7 and 9 and evaluated on runs 11 and 13, with 2 s windows at 0.5 s stride, each lying wholly inside one annotated period. The protocol alternates the two classes about every 4 s, a median of 5 windows per class block. Its analysis plan was committed to version control before any model was trained on it (Section 6.7).

**Cohort D (BCI Competition IV-2a, BNCI2014-001).** Nine participants recorded in two sessions on different days, 22-channel EEG at 250 Hz resampled to 100 Hz, cued motor imagery of the left against the right hand. The model is fitted on the first session and evaluated on the second, with 4 s windows covering the imagery interval, 288 trials per participant and balanced classes. Trial order is randomised, so a single-class block is one trial: this is the cohort where a covariance update spans the largest number of blocks. It contributes about 100 fitting windows per participant, the fewest of the five, which Section 6.10 shows breaks most of the published decoders under shared hyperparameters.

**Cohort E (lower-limb exoskeleton, DECODED).** Seven participants wearing a powered lower-limb exoskeleton, 27-channel EEG at 200 Hz resampled to 100 Hz, from the DECODED subproject of EUROBENCH. Each run presents 15 s standing relaxed with the device stopped, 24 s walking under kinesthetic motor imagery, 22 s walking under a regressive-counting distractor and 14 s standing relaxed. We decode standing relaxed against walking under motor imagery, the same contrast as cohort A on a second device, and drop the counting condition rather than folding it into either class, since it is a distractor rather than a movement-intent state. Windows are 4 s at 0.5 s stride and never cross a condition boundary. Only one participant recorded on two dates, so unlike cohorts A, C and D the split is temporal across runs, the first eight fitting and the rest testing. This cohort therefore matches best on task and tests least drift, and Section 8 says so.

The protocols differ sharply in temporal structure, which Section 6.6 shows to be decisive (Table 1).

**Table 1**

Contiguous single-class block structure, computed from the label sequence alone by `tools/tau_blk_rule.py`. Blocks are counted within a recording, never across one, and pooled over participants;  $\tau_{\text{blk}}$  is the median block duration in windows. It is a property of the protocol and not of the participant: every participant of every cohort returns the same median.

| Cohort                   | Blocks | $\tau_{\text{blk}}$ | Longest |
|--------------------------|--------|---------------------|---------|
| A (exoskeleton)          | 789    | 18                  | 486     |
| B (treadmill)            | 66     | 214                 | 2211    |
| C (motor execution)      | 3600   | 5                   | 5       |
| D (BCI IV-2a)            | 734    | 1                   | 6       |
| E (exoskeleton, DECODED) | 192    | 13                  | 33      |

## 5.2. Preprocessing and training

Signals are band-pass filtered to 8–30 Hz and resampled to 100 Hz. Cohort A uses a zero-phase FIR filter in MNE-Python [56], after which the EOG channels, filtered to the same band, are regressed linearly out of the EEG. Cohorts B and C use a zero-phase fourth-order Butterworth filter and have no EOG step. No other artefact rejection, ICA decomposition or channel interpolation is applied, so that the alignment layer is evaluated on the shift it is meant to correct. The band also excludes the gait cycle, near 1–2 Hz, and most electromyographic power above 30 Hz. Amplitude normalisation is applied globally per recording, before the squaring stage, for the reason given in Section 4.1.

Table 2 lists the training configuration. The adaptation momentum is  $m = 0.2$  unless stated otherwise; Section 6.6 treats  $m$  as the object of study. Model selection uses balanced accuracy on a held-out split of the fitting sessions, evaluated at the best epoch by that criterion, which is conventional but upward-biased on small folds (Section 8).

*Published decoders* Eight published decoders from the *braindecode* library [58] were trained in the same pipeline as the decoder: ShallowFBCSPNet and Deep4Net [10], EEG-Net [23], EEGNeX [24], EEG Conformer [27], TSception [59], FBLightConvNet and EEG-TCNet [28]. Each model's classification layer is replaced by the same linear projection and classifier, and each is trained with the preprocessing, optimiser, schedule, early stopping, model selection and temperature scaling of Table 2, so that the models differ only in the backbone. No per-model tuning was performed. Parameter counts include the shared classifier and are instantiated at cohort A's input shape.

## 5.3. Outcome measures

Balanced accuracy for comparability with the literature; expected calibration error (ECE) with 10 equal-width bins [35], because a device acting on a confidence value needs that value to mean something; accuracy at 90 % coverage, the trained selective operating point; spurious activations per minute of standing; and detection latency, reported alongside the benefit it buys.

**Table 2**

Training configuration, identical across all arms, all cohorts and the published decoders of Table 3, except where the cohort definition differs.

| Setting                     | Value                                                                           |
|-----------------------------|---------------------------------------------------------------------------------|
| Band-pass / resample        | 8–30 Hz, 100 Hz (zero-phase)                                                    |
| Artefact correction         | EOG regression (A); none (B, C)                                                 |
| Amplitude normalisation     | global, per recording                                                           |
| Window / stride             | 4 s / 0.5 s (A); 2 s (B); 2 s / 0.5 s (C)                                       |
| Context length $K$ , stride | 8 epochs, 3 (14.5 s span)                                                       |
| Optimiser                   | AdamW [57], lr $3 \times 10^{-4}$ , weight decay $10^{-2}$ , one-cycle schedule |
| Batch size $N$              | 32                                                                              |
| Epochs                      | at most 30; early stopping after 12 without improvement                         |
| Gradient clipping           | norm 1.0                                                                        |
| Loss                        | cross-entropy, inverse-frequency weighted                                       |
| Selective coverage target   | 0.90                                                                            |
| Adaptation momentum $m$     | 0.2 (swept in §6.6)                                                             |
| Model selection             | best held-out balanced accuracy                                                 |
| Calibration                 | temperature scaling fitted on the held-out split                                |

*Leakage guard* Unconditioned nuisance-decodability is confounded with accuracy. In our pipeline a decoder that uses no movement information whatever measures  $R^2 = +0.863$  against the movement signal, because both the decoder output and the nuisance are driven by the class label. Movement leakage is therefore measured conditional on the class label throughout, by centring the nuisance features within class before regression.

We calibrate the probe by injecting the recorded movement signal into the EEG, spatially mixed with a fixed per-participant gain matrix, at a known fraction  $f$  of signal power. The calibration uses the tangent-space representation. Accuracy rises monotonically with  $f$ : 0.763 at  $f = 0$ , then 0.819, 0.873, 0.892 and 0.913 at  $f = 0.001, 0.005, 0.01$  and  $0.05$ . The conditional probe rises with it, from  $-0.112$  at  $f = 0$  through  $+0.015$  at  $f = 0.005$  to  $+0.067$  at  $f = 0.02$ , where accuracy is 0.903. Reaching an accuracy near 0.90 through injected movement therefore moves the probe to about  $+0.07$ . The proposed *align + gate* decoder reaches 0.901 with a conditional probe of  $-0.068$ , and all eight cohort-A arms lie between  $-0.066$  and  $-0.082$ , on the uncontaminated side of the calibration curve. Two limits apply. The calibration was performed on a different representation from the decoder it is compared with, and the conditional probe cannot register artefact information that is constant within each class. The comparison is therefore evidence against movement-driven accuracy. Section 6.3 adds a second, independent test by removing artefact components and retraining.

## 5.4. Statistical treatment and measurement noise

Drift reductions are tested with a two-sided paired  $t$ -test over participants, pairing each participant's aligned distance against its raw distance. Accuracy comparisons are reported as the mean over participants with the between-subject standard deviation, and as paired Wilcoxon signed-rank tests over participants wherever per-participant results were retained (Section 8). Every  $\pm$  in a table is labelled in its caption as across participants or across data-split seeds. Standard deviations across participants are population values, and standard deviations across seeds are sample values.

Two ablation arms disable alignment entirely and therefore never invoke the running estimator. For them the momentum and the estimator version are inert, so every run is a true repeat and they measure the noise floor directly (Table 10).

The stem-only arm reproduces 0.7375 bit-for-bit across three runs; the arms containing the cross-epoch transformer do not, and the attention path is the only structural difference between them. We attribute the variation to non-deterministic reductions in scaled dot-product attention on GPU.

Run-to-run variation at a fixed split seed is a lower bound on the relevant noise. Across three data-split seeds the *align + gate* configuration spans 0.839 to 0.884, a range of 0.046 and a standard deviation of 0.026 (Table 10), several times the fixed-seed figure. This is the quantity that governs whether two models can be ranked, and differences between models smaller than it are reported without a ranking. The spread also differs between arms. On cohort B at a memory of 320 windows it is 0.003 for *align + gate* and 0.037 for the no-align arm, which contains the transformer, so single-seed values of the no-align arm on that cohort lie inside its own seed spread.

Any arm containing the transformer carries approximately 0.010 run-to-run noise at a fixed seed, so differences below 0.02 are not interpreted anywhere in this paper. The effects we do interpret on accuracy are measured over three data-split seeds and tested paired within participant (Tables 3 and 6, Section 6.4). An exact null in the deterministic arm is stronger evidence than one merely within noise, because that arm has no noise to hide in.

*Pipeline separation* The proposed decoder and the published decoders of Section 6.11 were trained in one pipeline. Tables 7 and 8 come from earlier pipelines and are used only for comparisons within each table.

## 6. Results

### 6.1. The drift mechanism


We measure the Riemannian distance (1) between the fitting-session mean covariance and each held-out session's, before and after the layer (Fig. 2, Table 9).

Drift is reduced in 15 of 15 participants ( $t = -9.46$ ,  $p = 8 \times 10^{-5}$  for cohort A;  $t = -12.33$ ,  $p = 5 \times 10^{-6}$  for cohort B), with per-participant reductions spanning 38.7 to

**Table 3**

The reported model, the stem it is built on, and eight published decoders trained in one pipeline with identical preprocessing, optimiser, schedule, early stopping, model selection and classifier head. Balanced accuracy, mean  $\pm$  SD across three data-split seeds. Cohorts A, B, C, D and E have 7, 8, 20, 9 and 7 participants. Bold marks a configuration of ours that exceeds every published decoder in that column. The branch leads four columns; on cohort B, where a class block outlasts one covariance update and Condition 1 rules the branch out, the stem alone leads instead, so one of the two is the best entry in every column and the protocol selects between them before training. Parameter counts include the shared classifier at cohort A's input shape. † near chance ( $< 0.55$ ) for at least five participants on a seed, a training failure under the shared settings. These columns are not comparable with numbers from the decoders' original papers.

| Model                                    | Params  | A                                   | B                                   | C                                   | D                                   | E                                   |
|------------------------------------------|---------|-------------------------------------|-------------------------------------|-------------------------------------|-------------------------------------|-------------------------------------|
| <i>This work</i>                         |         |                                     |                                     |                                     |                                     |                                     |
| stem + branch                            | 290 943 | <b>0.909 <math>\pm</math> 0.004</b> | 0.568 $\pm$ 0.005                   | <b>0.843 <math>\pm</math> 0.006</b> | <b>0.741 <math>\pm</math> 0.011</b> | <b>0.949 <math>\pm</math> 0.005</b> |
| stem alone                               | 19 505  | 0.864 $\pm$ 0.032                   | <b>0.744 <math>\pm</math> 0.016</b> | 0.779 $\pm$ 0.007                   | <b>0.701 <math>\pm</math> 0.023</b> | <b>0.928 <math>\pm</math> 0.006</b> |
| <i>Published decoders, same pipeline</i> |         |                                     |                                     |                                     |                                     |                                     |
| EEGNeX [24]                              | 71 010  | 0.876 $\pm$ 0.014                   | 0.707 $\pm$ 0.023                   | 0.783 $\pm$ 0.004                   | 0.656 $\pm$ 0.015                   | 0.887 $\pm$ 0.018                   |
| EEGConformer [27]                        | 445 442 | 0.847 $\pm$ 0.008                   | 0.685 $\pm$ 0.018                   | 0.783 $\pm$ 0.014                   | 0.599 $\pm$ 0.077†                  | 0.896 $\pm$ 0.010                   |
| FBLightConvNet                           | 23 760  | 0.866 $\pm$ 0.011                   | 0.687 $\pm$ 0.010                   | 0.778 $\pm$ 0.006                   | 0.674 $\pm$ 0.004                   | 0.891 $\pm$ 0.022                   |
| ShallowFBCSPNet [10]                     | 206 962 | 0.849 $\pm$ 0.020                   | 0.692 $\pm$ 0.027                   | 0.765 $\pm$ 0.004                   | 0.682 $\pm$ 0.013                   | 0.886 $\pm$ 0.008                   |
| TSception [59]                           | 182 098 | 0.852 $\pm$ 0.003                   | 0.734 $\pm$ 0.017                   | 0.723 $\pm$ 0.001                   | 0.579 $\pm$ 0.015†                  | 0.833 $\pm$ 0.011                   |
| EEGNet [23]                              | 27 666  | 0.856 $\pm$ 0.028                   | 0.716 $\pm$ 0.025                   | 0.630 $\pm$ 0.049†                  | 0.526 $\pm$ 0.017†                  | 0.700 $\pm$ 0.054                   |
| EEGTCNet [28]                            | 7 062   | 0.878 $\pm$ 0.009                   | 0.672 $\pm$ 0.024                   | 0.585 $\pm$ 0.037†                  | 0.505 $\pm$ 0.015†                  | 0.695 $\pm$ 0.016                   |
| Deep4Net [10]                            | 717 417 | 0.817 $\pm$ 0.014                   | 0.617 $\pm$ 0.030                   | 0.549 $\pm$ 0.004†                  | 0.525 $\pm$ 0.007†                  | 0.722 $\pm$ 0.035                   |



**Figure 2:** Session-drift reduction per participant. Bars are the Riemannian distance  $\delta$  between the fitting-session mean covariance and the held-out session mean, for the raw signal, the momentum-0 no-op control, and the aligned signal. (a) cohort A,  $n = 7$ ; (b) cohort B,  $n = 8$ . Distance falls in every participant of both cohorts; the no-op control, which shares every code path but never updates its estimate, leaves the distance unchanged.

**Table 4**

Component ablation on cohort A under the current estimator: balanced accuracy (mean  $\pm$  SD across seeds), ECE, accuracy at 90% coverage and spurious activations per minute of standing, each averaged over seeds. Acc@90 is withheld for arms without a trained gate, whose confidence ranks windows in recording order. Seed 0 of the stem-only arm predates the estimator fix and is excluded. The last three columns compare each arm with the stem-only arm over seeds 1 and 2, which every arm shares, paired over the 7 participants: difference in mean accuracy, participants higher than the stem-only arm, and the Wilcoxon signed-rank  $p$ , Holm-corrected over the six contrasts. None of the arms in this table is the reported model, which is the stem plus the tangent branch; these are the components of Section 4.3, and the reported model is compared with the same decoders in Table 3.

| Components         | Seeds  | Acc               | ECE   | Acc@90 | FA/min | $\Delta$ stem | Higher | $p_{\text{Holm}}$ |
|--------------------|--------|-------------------|-------|--------|--------|---------------|--------|-------------------|
| align + gate       | 3 of 3 | 0.868 $\pm$ 0.026 | 0.058 | 0.835  | 1.60   | -0.022        | 3 of 7 | 1.00              |
| align only         | 3 of 3 | 0.869 $\pm$ 0.017 | 0.054 | —      | 1.54   | -0.020        | 3 of 7 | 1.00              |
| gate only          | 3 of 3 | 0.842 $\pm$ 0.022 | 0.076 | 0.827  | 0.89   | -0.027        | 2 of 7 | 0.94              |
| align + ctx + gate | 3 of 3 | 0.859 $\pm$ 0.002 | 0.072 | 0.818  | 1.16   | -0.025        | 3 of 7 | 1.00              |
| align + ctx        | 3 of 3 | 0.857 $\pm$ 0.004 | 0.074 | —      | 1.00   | -0.027        | 1 of 7 | 1.00              |
| ctx + gate         | 3 of 3 | 0.854 $\pm$ 0.017 | 0.080 | 0.796  | 0.66   | -0.019        | 2 of 7 | 1.00              |
| stem only          | 3 of 3 | 0.864 $\pm$ 0.032 | 0.063 | —      | 0.70   | —             | —      | —                 |

51.2% on cohort A and 43.8 to 82.1% on cohort B (Table 9). No participant regresses, and the reduction is not carried by one or two outliers. The no-op control leaves the distance unchanged, which establishes that the effect comes from adaptation and excludes both the whitening operation and the measurement itself.

The no-op control runs at momentum 0. Its running estimate is primed by copying the first batch's covariance and never updates, so it applies one fixed linear map to every session,  $M = gW + (1 - g)I$  with  $g = \sigma(2.0) \approx 0.88$ . A fixed map applied to both covariances is a congruence, under which the affine-invariant distance (1) does not change, and the control accordingly changes the distance by less than 0.1% of its raw value in every participant (Table 9). An earlier version of this measurement normalised each covariance by its trace after the map, which breaks that invariance and made the control scatter in both directions; it was removed, and the values reported here replace the larger reductions it produced (52.0 and 83.9%).

The mechanism works as designed, and Section 6.4 shows that this does not settle whether it helps the task.

## 6.2. Component ablation, cohort A

Table 12 places the reported model in its own ablation. The branch is the only component that raises accuracy over the stem on this cohort. Of the configurations built from the alignment layer, the gate and the cross-epoch transformer, none improves on the stem by more than 0.005, and the gate alone is 0.022 below it, the one difference in the table that reaches significance in the negative direction. Adding alignment on top of the branch costs 0.032 relative to the branch alone, so the two are substitutes rather than complements.

Note also the SD column, which is across seeds rather than participants. The reported model is the most reproducible arm in the table. That matters because the interpretation threshold this project uses was set by the seed spread

of the transformer arms, and the branch narrows the very quantity that set it.

Table 4 repeats the ablation under the current estimator, over three data-split seeds for six arms and two for the stem-only arm, and compares every arm with the stem-only arm paired over participants. Where it and the single-seed table disagree, the paragraphs below follow Table 4.

The acc@90 value for the *align + ctx* arm is withheld because with the selection head disabled the model emits near-constant confidence, so ranking by it selects the first 90% of the test set in recording order. That number would be an ordering artefact.

**Selective prediction** Abstention is supposed to raise accuracy on the windows that are kept. It does not. Across the 18 arms in this work for which both quantities were computed on the same windows with the same metric, declining the least-confident 10% **lowers** balanced accuracy in 14, by 0.045 on average; the headline arm loses 0.026 (0.901  $\rightarrow$  0.876) and the worst case loses 0.122. The four arms where abstention helps are exactly those with no trained gate (*align + ctx*, *stem only*), where the ranking is arbitrary. A gate that is beaten by its own absence is not ranking windows by correctness, and the selective head is therefore reported as a component for which no benefit was demonstrated. Fixing it, for instance by thresholding a temperature-scaled softmax confidence instead of the learned gate, is left open. Calibration is compared with the published decoders in Section 6.11.


The head contributes no accuracy either. Under the trace-normalised estimator, an *align*-only arm reaches 0.883 against 0.884 with the head, a difference of 0.002, and adding the head to *align + ctx* changes accuracy by 0.000 (0.862 in both arms). Calibration does not suffer either: without the head ECE is 0.038, against 0.042 with it, both under the same estimator. Over three data-split seeds the *align*-only arm averages 0.869 against 0.868 with the head, a difference of -0.001 (higher with the head in 4 of 7 participants, Wilcoxon  $p = 1.00$ ), and mean ECE is 0.054 without the head against 0.058 with it. An earlier difference of +0.016 between the same two context arms under the original estimator is 1.6 $\times$  the 0.010 noise floor of arms containing the transformer (Table 10), and it does not reappear in the comparison made with one estimator throughout. On this evidence the selective head could be removed with no loss of accuracy or calibration. It is kept in the arms reported here so that they remain comparable with the ablation.

**Dependence on the alignment stage** With alignment removed, the context and selection modules together contribute between 0.000 and 0.007: the no-align arm spans 0.827–0.834 across repeats (Table 10) while the stem-only arm is deterministic at 0.827. Against the 0.010 noise floor that is indistinguishable from nothing, and well below alignment's own contribution: +0.051 under the original estimator (0.878 against 0.827, Table 11) and +0.028 under the trace-normalised one (0.862 against 0.834) from the

**Table 5**

The adaptation-rate condition, single seed. Accuracies are mean  $\pm$  SD across participants. Momentum  $m$  sets the estimator's memory to roughly  $32/m$  windows: 160, 640 and 3200. **Every arm that uses alignment gains as the memory crosses the class-block length, and neither arm that omits it responds**; the stem-only arm reproduces 0.737 exactly at all three settings despite a 160-fold change. The lower block applies the identical change to the cohort whose blocks are short, where it costs 0.060, so the two cohorts have opposite optima. Three-seed values for cohort B are in Table 6, and for cohort A in Section 6.4. None of the arms in this table is the reported model, which is the stem plus the tangent branch; these are the components of Section 4.3, and the reported model is compared with the same decoders in Table 3.

| Arm                                                                                             | $m=0.20$          | $m=0.05$          | $m=0.01$          | $\Delta$ at crossing | Aligns    | $\Delta$ cohort A |
|-------------------------------------------------------------------------------------------------|-------------------|-------------------|-------------------|----------------------|-----------|-------------------|
| <i>Cohort B (<math>\tau_{\text{blk}} = 309</math>): slowing adaptation past the class block</i> |                   |                   |                   |                      |           |                   |
| align + ctx + gate                                                                              | $0.631 \pm 0.092$ | $0.724 \pm 0.097$ | $0.723 \pm 0.120$ | <b>+0.093</b>        | yes       | —                 |
| align + gate                                                                                    | $0.581 \pm 0.044$ | $0.671 \pm 0.063$ | $0.696 \pm 0.089$ | <b>+0.090</b>        | yes       | —                 |
| align + ctx                                                                                     | $0.606 \pm 0.063$ | $0.700 \pm 0.082$ | $0.726 \pm 0.070$ | <b>+0.094</b>        | yes       | —                 |
| ctx + gate                                                                                      | $0.788 \pm 0.097$ | $0.782 \pm 0.111$ | $0.783 \pm 0.105$ | -0.007               | <b>no</b> | —                 |
| stem only                                                                                       | $0.737 \pm 0.104$ | $0.737 \pm 0.104$ | $0.737 \pm 0.104$ | +0.000               | <b>no</b> | —                 |
| <i>Cohort A (<math>\tau_{\text{blk}} = 18</math>): the same change, matched estimator</i>       |                   |                   |                   |                      |           |                   |
| align + ctx + gate                                                                              | $0.862 \pm 0.145$ | —                 | $0.801 \pm 0.107$ | <b>-0.060</b>        | yes       | <b>-0.060</b>     |
| ctx + gate                                                                                      | $0.834 \pm 0.094$ | —                 | $0.827 \pm 0.116$ | -0.007               | <b>no</b> | -0.007            |



**Figure 3:** Component ablation on cohort A, seed 0, over the components of Section 4.3 that were tested and are not in the reported model. It is retained because the ablation is reported against these arms; the reported model appears in Table 12 and Fig. 8. (a) Balanced accuracy by arm; error bars are  $\pm 1$  SD across the 7 participants; seed variation is reported separately in Table 10. (b) Balanced accuracy at 90% coverage, obtained by declining the least-confident 10% of windows. It sits *below* full-coverage accuracy for every arm with a trained gate. (c) Balanced accuracy against spurious activations per minute of standing, showing the cross-epoch context as a selectable operating point. Three-seed values under the current estimator are in Table 4.

**Table 6**

Cohort B over three data-split seeds under the current estimator. *Cost* is *align + gate* minus the gate-only arm ( $0.729 \pm 0.006$ ), which differs only in alignment and does not use the estimator, paired within participant with seeds averaged; *p* is the Wilcoxon signed-rank test. The cost falls gradually as the memory lengthens, with no step at  $\tau_{\text{blk}} = 309$ , and it is no longer significant at memories of 1600 and 3200 windows. None of the arms in this table is the reported model, which is the stem plus the tangent branch; these are the components of Section 4.3, and the reported model is compared with the same decoders in Table 3.

| <i>m</i> | Memory | align + gate      | Cost   | Lower in | <i>p</i> |
|----------|--------|-------------------|--------|----------|----------|
| 0.20     | 160    | $0.580 \pm 0.008$ | -0.149 | 8 of 8   | 0.008    |
| 0.10     | 320    | $0.638 \pm 0.002$ | -0.091 | 7 of 8   | 0.016    |
| 0.02     | 1600   | $0.695 \pm 0.005$ | -0.034 | 6 of 8   | 0.148    |
| 0.01     | 3200   | $0.698 \pm 0.002$ | -0.031 | 6 of 8   | 0.148    |

**Table 7**

Two-condition comparisons. **The first block is the identifying evidence for the stem design:** removing per-window normalisation is worth +0.134 and +0.174 where the representation is pure marginal power, and **exactly +0.000 where it is scale-free by construction** and therefore cannot respond. An effect that scales with power content and vanishes without it identifies the mechanism instead of merely demonstrating it. The second block retired an amplitude side-channel: supplying a per-window amplitude feature and supplying the same feature with its window assignment shuffled differ by under 0.02 for four of five models. Single seed; condition and control accuracies are mean  $\pm$  SD across participants.

| Comparison                                                                                       | Power content | Condition         | Control           | $\Delta$      |
|--------------------------------------------------------------------------------------------------|---------------|-------------------|-------------------|---------------|
| <i>Per-window amplitude normalisation, by how much marginal power the representation carries</i> |               |                   |                   |               |
| band power, cohort A                                                                             | full          | $0.625 \pm 0.091$ | $0.760 \pm 0.113$ | +0.134        |
| band power, cohort B                                                                             | full          | $0.574 \pm 0.084$ | $0.748 \pm 0.067$ | +0.174        |
| tangent space, cohort A                                                                          | partial       | $0.762 \pm 0.146$ | $0.746 \pm 0.138$ | -0.016        |
| tangent space, cohort B                                                                          | partial       | $0.675 \pm 0.088$ | $0.704 \pm 0.085$ | +0.029        |
| correlation, cohort A                                                                            | <b>none</b>   | $0.738 \pm 0.138$ | $0.734 \pm 0.134$ | <b>-0.004</b> |
| correlation, cohort B                                                                            | <b>none</b>   | $0.701 \pm 0.058$ | $0.702 \pm 0.058$ | <b>+0.000</b> |
| <i>Amplitude side-channel against its shuffled control</i>                                       |               |                   |                   |               |
| ShallowFBCSPNet                                                                                  | —             | $0.832 \pm 0.078$ | $0.816 \pm 0.101$ | -0.017        |
| Deep4Net                                                                                         | —             | $0.798 \pm 0.142$ | $0.809 \pm 0.123$ | +0.011        |
| EEGNeX                                                                                           | —             | $0.841 \pm 0.093$ | $0.841 \pm 0.086$ | -0.000        |
| EEGITNet                                                                                         | —             | $0.828 \pm 0.115$ | $0.837 \pm 0.111$ | +0.008        |
| BDTCN                                                                                            | —             | $0.748 \pm 0.136$ | $0.652 \pm 0.151$ | -0.095        |

estimator-fix run of Section 6.5. The lower figure should be read with care: the no-align arm is inert to the estimator, so its 0.834 differs from 0.827 only by run-to-run variation, and part of the gap between the two figures is that noise. The comparison that isolates alignment most directly, *align + gate* against *gate* with nothing else changed, was repeated over three data-split seeds under the trace-normalised estimator. With each participant's accuracy averaged over seeds, alignment raises accuracy by +0.026 (0.868 against 0.842), higher in 6 of 7 participants (Wilcoxon  $p = 0.22$ ). The effect is +0.067, +0.029 and -0.019 on the three seeds, and it is uneven across participants: sub-04 gains 0.179 and sub-05 loses 0.127, each with the same sign on every seed. The gain from alignment on cohort A is therefore close to the trace-normalised single-seed figure of +0.028, and with seven participants it is not significant. Table 4 removes the

**Table 8**

Where amplitude normalisation is applied, across 6 published decoders on cohort A, single seed, in an earlier training pipeline. All three columns are balanced accuracies, mean  $\pm$  SD across participants; best mean per row in bold. Per-window normalisation is best for 2 and worst for 2, and most differences are below the 0.02 noise threshold, so this comparison neither supports nor contradicts the representation-level result in Table 7.

| Model                | Per-window                          | Global                              | None                                |
|----------------------|-------------------------------------|-------------------------------------|-------------------------------------|
| ShallowFBCSPNet [10] | <b><math>0.824 \pm 0.098</math></b> | $0.821 \pm 0.087$                   | $0.818 \pm 0.096$                   |
| Deep4Net [10]        | $0.760 \pm 0.090$                   | <b><math>0.821 \pm 0.131</math></b> | $0.821 \pm 0.121$                   |
| EEGNeX [24]          | $0.834 \pm 0.085$                   | $0.832 \pm 0.104$                   | <b><math>0.835 \pm 0.097</math></b> |
| EEGITNet [60]        | <b><math>0.833 \pm 0.096</math></b> | $0.791 \pm 0.134$                   | $0.795 \pm 0.137$                   |
| BDTCN                | $0.820 \pm 0.048$                   | <b><math>0.822 \pm 0.080</math></b> | $0.820 \pm 0.056$                   |
| EEGConformer [27]    | $0.828 \pm 0.066$                   | <b><math>0.829 \pm 0.079</math></b> | $0.805 \pm 0.074$                   |

**Table 9**

Per-participant session drift.  $\delta$  is the affine-invariant Riemannian distance between the fitting-session mean covariance and the held-out mean. The no-op control shares every code path but never updates its estimate. **Every participant improves:** 15 of 15,  $t = -9.46$ ,  $p = 8 \times 10^{-5}$  (A) and  $t = -12.33$ ,  $p = 5 \times 10^{-6}$  (B). The control changes the distance by at most 0.09 % of its raw value. The row under each cohort is the mean  $\pm$  SD across participants.

| Participant                          | Raw $\delta$      | No-op             | Aligned           | Reduction         |
|--------------------------------------|-------------------|-------------------|-------------------|-------------------|
| <i>Cohort A, 6 held-out sessions</i> |                   |                   |                   |                   |
| A01                                  | 6.925             | 6.926             | 3.385             | 51.1 %            |
| A02                                  | 11.867            | 11.878            | 7.273             | 38.7 %            |
| A03                                  | 5.891             | 5.891             | 2.875             | 51.2 %            |
| A04                                  | 10.337            | 10.339            | 5.240             | 49.3 %            |
| A05                                  | 5.208             | 5.208             | 3.163             | 39.3 %            |
| A06                                  | 6.909             | 6.909             | 3.650             | 47.2 %            |
| A07                                  | 8.986             | 8.987             | 4.986             | 44.5 %            |
| Mean $\pm$ SD                        | $8.017 \pm 2.267$ | $8.020 \pm 2.270$ | $4.367 \pm 1.451$ | $45.9 \pm 4.9$ %  |
| <i>Cohort B, 2 held-out trials</i>   |                   |                   |                   |                   |
| SL01                                 | 6.658             | 6.658             | 3.176             | 52.3 %            |
| SL02                                 | 8.122             | 8.126             | 3.623             | 55.4 %            |
| SL03                                 | 5.739             | 5.739             | 1.879             | 67.3 %            |
| SL04                                 | 4.803             | 4.803             | 2.701             | 43.8 %            |
| SL05                                 | 6.153             | 6.154             | 2.050             | 66.7 %            |
| SL06                                 | 6.696             | 6.696             | 1.546             | 76.9 %            |
| SL07                                 | 4.737             | 4.737             | 0.846             | 82.1 %            |
| SL08                                 | 6.462             | 6.462             | 2.854             | 55.8 %            |
| Mean $\pm$ SD                        | $6.171 \pm 1.032$ | $6.172 \pm 1.033$ | $2.334 \pm 0.859$ | $62.5 \pm 12.2$ % |

single-seed ordering as well. Over the two seeds every arm shares, the stem-only arm leads: every other arm lies 0.019 to 0.027 below it paired over those seeds and is higher in at most 3 of 7 participants, and no difference is distinguishable from zero (Holm-corrected  $p \geq 0.94$ ). Over three seeds the stem-only arm averages  $0.864 \pm 0.032$ , level with the 0.868 of *align + gate*, because its seed 0 sits at 0.827 against 0.881 and 0.884 on the other two. The ranking between these arms is therefore set by the seed, not by the components. Against the full configuration the stem's lead is carried by two participants (sub-05 +0.137, sub-01 +0.075) and reversed by a third (sub-04 -0.130). The alignment arms also average more spurious activations per minute than the corresponding arms without alignment (1.60 against 0.89 over three seeds with the gate, and 1.17 against 0.74 over seeds 1 and 2 with context and gate), yet they are higher in only 3 of 7 participants ( $p = 0.44$  and  $p = 0.81$ ). On cohort

**Table 10**

Measurement noise. The first block repeats arms for which momentum and estimator version are inert, so all variation is run-to-run: the arm without the cross-epoch transformer is bit-identical across three runs, and the arms with it are not. Across data-split seeds the spread is larger again, **0.032 to 0.046** for our arms without the transformer. **Differences below 0.02 are not interpreted anywhere in this work**, and differences below the seed spread are reported without a ranking. None of the arms in this table is the reported model, which is the stem plus the tangent branch; these are the components of Section 4.3, and the reported model is compared with the same decoders in Table 3.

| Arm                                                                              | Transformer | <i>n</i> | Spread        |
|----------------------------------------------------------------------------------|-------------|----------|---------------|
| <i>Fixed seed, repeated runs of arms where the manipulated variable is inert</i> |             |          |               |
| stem only, cohort B                                                              | no          | 3        | <b>0.0000</b> |
| no-align, cohort A                                                               | yes         | 3        | 0.0069        |
| no-align, cohort B                                                               | yes         | 4        | 0.0107        |
| <i>Across three data-split seeds, trace-normalised estimator</i>                 |             |          |               |
| align + gate                                                                     | no          | 3        | 0.0455        |
| gate only                                                                        | no          | 3        | 0.0409        |
| align only                                                                       | no          | 3        | 0.0320        |
| <i>Cohort B, memory 320 windows, three data-split seeds</i>                      |             |          |               |
| align + gate                                                                     | no          | 3        | 0.0028        |
| no-align                                                                         | yes         | 3        | 0.0369        |
| <i>Across data-split seeds, published decoders, same pipeline</i>                |             |          |               |
| Deep4Net                                                                         | —           | 3        | 0.0265        |
| EEGConformer                                                                     | —           | 3        | 0.0153        |
| EEGNeX                                                                           | —           | 3        | 0.0268        |
| EEGNet                                                                           | —           | 3        | 0.0518        |
| EEGTCNet                                                                         | —           | 3        | 0.0175        |
| FBLightConvNet                                                                   | —           | 3        | 0.0215        |
| ShallowFBCSPNet                                                                  | —           | 3        | 0.0393        |
| TSception                                                                        | —           | 3        | 0.0062        |

**Table 11**

Component ablation, cohort A, single seed. Accuracy is the mean over 7 participants with the between-subject SD. Acc@90 is accuracy at 90 % coverage; FA/min is spurious activations per minute of standing. All cells use the original estimator so that arms are mutually comparable. Three-seed values under the current estimator are in Table 4. The arms here are the components of Section 4.3, and none of them is the reported model; that is the stem plus the tangent branch (Table 3).

| Components         | Acc $\pm$ SD                        | ECE          | Acc@90       | FA/min |
|--------------------|-------------------------------------|--------------|--------------|--------|
| align + gate       | <b>0.901 <math>\pm</math> 0.059</b> | <b>0.039</b> | <b>0.876</b> | 1.64   |
| align + ctx + gate | 0.878 $\pm$ 0.102                   | 0.074        | 0.837        | 1.26   |
| align + ctx        | 0.862 $\pm$ 0.125                   | 0.072        | —            | 1.26   |
| ctx + gate         | 0.827 $\pm$ 0.096                   | 0.091        | 0.778        | 0.73   |
| stem only          | 0.827 $\pm$ 0.124                   | 0.082        | 0.823        | 0.73   |

A, therefore, neither alignment nor the other modules are shown to raise accuracy over the stem, and none is shown to lower it. The measured effect of alignment on this cohort is the reduction in session drift (Section 6.1).

**Cross-epoch context** Over three seeds under the current estimator (Table 4), adding it changes accuracy from 0.868 to 0.859, a difference below the 0.02 this paper interprets for arms containing the transformer, raises ECE from 0.058 to 0.072, and reduces spurious activations from 1.60 to 1.16 per

**Table 12**

Ablation on cohort A containing the reported model, three data-split seeds where available and two otherwise, under the current estimator. Accuracy is the grand mean over the 7 participants and SD is across seeds, not across participants, so it measures reproducibility rather than spread.  $\Delta$  and the test are paired against the stem alone over participants with seeds averaged within participant. The reported model is the first row. Adding the alignment layer to it (second row) removes most of the gain, which is the substitution reported in Section 6.9: both read the same second-order structure.

| Components                   | Seeds | Acc          | SD <sub>seed</sub> | ECE   | $\Delta$ vs stem | <i>p</i> (higher) |
|------------------------------|-------|--------------|--------------------|-------|------------------|-------------------|
| stem + branch                | 3     | <b>0.909</b> | 0.004              | 0.049 | +0.045           | 0.219 (6/7)       |
| stem + align + gate + branch | 3     | 0.877        | 0.018              | 0.069 | +0.013           | 0.578 (5/7)       |
| stem + align + gate          | 3     | 0.868        | 0.026              | 0.058 | +0.004           | 0.938 (4/7)       |
| stem + align                 | 2     | 0.863        | 0.017              | 0.062 | −0.001           | 0.812 (3/7)       |
| stem + gate                  | 3     | 0.842        | 0.022              | 0.076 | −0.022           | 0.047 (1/7)       |
| stem + gate + ctx            | 2     | 0.864        | 0.000              | 0.077 | −0.000           | 0.938 (4/7)       |
| stem + align + ctx           | 2     | 0.855        | 0.000              | 0.079 | −0.009           | 0.688 (3/7)       |
| stem + align + gate + ctx    | 2     | 0.858        | 0.001              | 0.077 | −0.006           | 0.812 (3/7)       |
| stem alone                   | 3     | 0.864        | 0.032              | 0.063 | —                | —                 |

**Table 13**


Artefact control on cohort A. Balanced accuracy of the proposed configuration on the training task, before and after ICA cleaning, with the number of muscle components removed per recording. The last row is the mean  $\pm$  SD across participants.

| Participant   | Uncleaned         | Cleaned           | Change             | Muscle removed |
|---------------|-------------------|-------------------|--------------------|----------------|
| sub-01        | 0.969             | 0.955             | −0.014             | 2.9            |
| sub-02        | 0.898             | 0.914             | +0.016             | 7.1            |
| sub-03        | 0.955             | 0.841             | −0.114             | 2.3            |
| sub-04        | 0.835             | 0.817             | −0.018             | 1.9            |
| sub-05        | 0.839             | 0.893             | +0.053             | 1.3            |
| sub-06        | 0.673             | 0.548             | −0.124             | 0.3            |
| sub-07        | 0.680             | 0.525             | −0.155             | 2.3            |
| Mean $\pm$ SD | 0.835 $\pm$ 0.111 | 0.785 $\pm$ 0.163 | −0.051 $\pm$ 0.074 |                |

minute, a 27 % reduction. Paired tests are not available because seed 0 of the context arm did not retain per-participant rows. It is therefore reported as a selectable safety mode, outside the default configuration. Removing it also removes 594,816 of the model's 618,997 parameters, which is what makes the default configuration compact.

### 6.3. Artefact control

Walk and stop differ in movement, so accuracy on this task could arise from muscle or motion artefact. The proposed configuration was therefore retrained after independent component analysis, with both arms restricted to the training task so that they differ only in the cleaning step. Each recording was re-referenced to the average, a copy band-passed to 1–45 Hz was decomposed into 30 components with the extended Picard algorithm [61], and the components that ICLabel [62] classified as eye, muscle, heart, line noise or channel noise were removed before band-passing to the decoding band. Across the 63 recordings, a mean of 6.8 of 30 components was removed, most often eye (42 % of removals) and muscle (38 %). Gait-related artefact is difficult to separate from cortical activity during walking [63, 64], and ICLabel is not trained on it specifically, so this control narrows the question and leaves part of it open.



**Figure 4:** Scalp maps for one cohort A participant (sub-01), fitting sessions. (a) Relative channel power in the layer's trace-normalised covariance estimate. (b) Per-channel gain of the whitening transform  $M^{-1/2}$ . (c) Absolute 8–30 Hz power during walking relative to standing, in dB. Power falls over sensorimotor electrodes, the direction expected for mu and beta event-related desynchronisation, and rises over occipital and parieto-occipital electrodes, where neck muscle activity and visual input during walking would both appear.

Cleaning lowers balanced accuracy from 0.835 to 0.785, a paired difference of  $-0.051$  (higher in 2 participants, lower in 5; Wilcoxon  $p = 0.22$ ; Table 13). Most of the accuracy therefore survives, well above the 0.500 chance level, and four participants are essentially unchanged or improve. The loss is concentrated in three participants, of 0.11 to 0.16. It does not follow the amount of muscle activity removed. The participant with the most muscle components removed, 7.1 per recording, improves by 0.016, and one with 0.3 removed loses 0.124 (Spearman  $\rho = +0.22$  between muscle components removed and accuracy change). Dependence on muscle artefact would produce the opposite pattern. Two of the three largest losses fall on the participants with the lowest baseline accuracy, 0.673 and 0.680, where estimates are noisiest. The conditional leakage probe stays negative in both arms ( $-0.088$  and  $-0.077$ ). With seven participants, and with re-referencing also part of the cleaning, the control falls short of establishing that the accuracy is neural, and it gives no support to the view that the accuracy rests on muscle artefact.

Fig. 4 shows the scalp distribution for one participant. Absolute 8–30 Hz power falls over sensorimotor electrodes during walking, as mu and beta desynchronisation predicts [46, 47], and rises over occipital and parieto-occipital electrodes, where neck muscle activity and changing visual input during walking would both appear. With one participant this shows only that the two effects are spatially separate. It does not show which of them the decoder relies on, and the retraining control above remains the test of that.

#### 6.4. What the alignment layer did, and why it is not retained

The alignment layer is not in the reported model, and the measurements that removed it are worth one section because they are the same phenomenon the branch shows.


*The ordering inverts between cohorts* On cohort A the arms that use alignment lead and the alignment-free arms trail; on cohort B the order reverses completely (Table 14, Fig. 5). With alignment enabled at the default rate the configuration averages 0.580 on cohort B, below every published decoder in Table 3, while the alignment-free arm reaches 0.792. This is not a transfer failure of a particular architecture. It is the same reversal the branch shows on the same cohort, for the same reason: both mechanisms divide by a running covariance, and cohort B is the cohort whose class blocks outlast the estimator memory, so that estimate converges on the streaming class (Section 6.9).

*Slowing the estimator has a cost where blocks are short* Lengthening the estimator memory from 160 to 3200 windows recovers cohort B, and costs accuracy on both cohorts whose blocks are shorter than one update:  $-0.050$  on cohort A and  $-0.016$  on cohort C, the latter lower in 15 of 20 participants ( $p = 0.003$ ) and negative on all three seeds ( $-0.085$ ,  $-0.056$ ,  $-0.008$ ). A rate cannot therefore be chosen once for all protocols, which is what makes the condition two-sided in principle even though only its lower side is measured here (Section 8).

*Why it is not retained* Added to the stem, the layer changes accuracy by  $+0.004$  on cohort A, inside the interpretation threshold (Table 12). Added on top of the branch it costs 0.032, and on cohort C the same addition is worth  $+0.002$  ( $p = 0.52$ ). Both mechanisms read the same second-order structure, so they are substitutes, and the branch is the one that raises accuracy.

#### 6.5. Correcting the covariance estimator

One candidate explanation is the power-weighted estimator described in Section 4.3. The defect is real and measurable: on cohort B, where the majority class carries  $1.52\times$  the power, the un-normalised running estimate sat at



**Figure 5:** The same five arms on both cohorts, single seed. None of these arms is the reported model; they are the components of Section 4.3. Solid heavy black is the best of them on cohort A; the other alignment-using arms are dashed in colour, and the two alignment-free arms are dashed and dotted. The three alignment arms cross below both alignment-free arms, so the ordering inverts instead of compressing.

**Table 14**

Both cohorts, single seed, original estimator. Balanced accuracy, mean  $\pm$  SD across participants. Three-seed cohort B values for *align + gate* are in Table 6. The arms here are the components of Section 4.3, and none of them is the reported model; that is the stem plus the tangent branch (Table 3).

| Components         | Cohort A                 | Cohort B                 |
|--------------------|--------------------------|--------------------------|
| align + gate       | <b>0.901</b> $\pm$ 0.059 | 0.581 $\pm$ 0.044        |
| align + ctx + gate | 0.878 $\pm$ 0.102        | 0.626 $\pm$ 0.073        |
| align + ctx        | 0.862 $\pm$ 0.125        | 0.606 $\pm$ 0.063        |
| ctx + gate         | 0.827 $\pm$ 0.096        | <b>0.792</b> $\pm$ 0.097 |
| stem only          | 0.827 $\pm$ 0.124        | 0.737 $\pm$ 0.104        |

Riemannian distance 0.306 from the walk covariance, which is to say it had become the majority class. The per-window trace normalisation in (4) moves it to 3.899 and 3.194 from the two class covariances.

Correcting it changed decoding by +0.005 on cohort B and -0.016 on cohort A. The estimator improved and the task did not follow, so this is therefore not the explanation for the inversion, though the correction is retained, which is independently right. It is the first of five cases in this work where a distribution-level quantity improved without the task following (Section 7).

## 6.6. The adaptation rate condition

**Condition 1 (Admissible rate band).** Assume (i) the running estimate is an exponential moving average with effective memory  $\mu$  windows; (ii) the within-class covariance is stationary over a block; (iii) the protocol presents classes in contiguous blocks of length  $\tau_{\text{blk}}$ ; and (iv) session drift acts as a congruence transform slow on the scale  $\tau_{\text{drift}}$ . Under

(i)–(iv) the estimate both avoids tracking the streaming class and tracks session drift only if

$$\tau_{\text{blk}} < \mu < \tau_{\text{drift}}. \quad (5)$$

Where  $\tau_{\text{blk}} \geq \tau_{\text{drift}}$  the band is empty and no setting of  $\mu$  satisfies both requirements. It is a design condition. Its lower bound is tested here empirically, once in advance; its upper bound is not measured, and we give no proof.

The running estimate tracks whatever is currently streaming. An exponential moving average with momentum  $m$  has an effective memory of approximately  $1/m$  batches, or  $NB/m$  windows at batch size  $N$  and  $B$  batches. When a protocol presents a class in contiguous blocks longer than that memory, the estimate converges to the currently streaming class, and whitening by it removes the between-class difference. Cohort A's blocks ( $\tau_{\text{blk}} = 18$ ) sit far inside even the fastest memory; cohort B's ( $\tau_{\text{blk}} = 214$ ) do not.


Fig. 6 shows the estimate itself for one participant from each cohort. It is the layer's internal state, which depends only on the input windows and the momentum and on no trained weight, so it is computed without training. At  $m = 0.2$  the estimate on cohort A stays away from both class covariances, while on cohort B it moves to the covariance of the class being streamed and follows it through the long walk block. At  $m = 0.01$  the estimate on cohort B still moves toward the walk covariance, more slowly, and crosses over late in the recording.

If this mechanism is responsible, slowing adaptation until the memory exceeds a class block should restore cohort B. The memory is approximately 160, 640 and 3200 windows at  $m = 0.2, 0.05$  and  $0.01$ , so the crossing of  $\tau_{\text{blk}} = 214$  falls between the first and second settings. Three seeds under the current estimator trace the recovery more finely for *align + gate* (Table 6), against the gate-only arm, which differs only in alignment and does not use the estimator (0.729). Alignment costs 0.149 at a memory of 160 windows (lower in 8 of 8 participants; Wilcoxon  $p = 0.008$ ), 0.091 at 320 (7 of 8;  $p = 0.016$ ), 0.034 at 1600 (6 of 8;  $p = 0.15$ ) and 0.031 at 3200 (6 of 8;  $p = 0.15$ ). The cost falls by similar amounts on either side of the block length and then levels off, so the recovery is gradual and shows no step at  $\tau_{\text{blk}} = 214$ . Alignment never becomes beneficial, and at the two slowest rates its remaining cost is not significant.


At seed 0, every arm that uses alignment gains +0.090 to +0.098 between memories of 160 and 640 windows (Table 5). Neither arm that omits alignment responds: the no-align arm moves -0.010, within its 0.0107 run-to-run spread (Table 10), and the stem-only arm reproduces 0.7375 exactly at every setting despite a 160-fold change in the adaptation rate. Because that arm is deterministic, this null is exact. The manipulated variable therefore acts only through the component it controls, which an unrelated confound such as optimisation dynamics or regularisation would not produce.

## 6.7. A prediction registered before a third cohort

Condition 1 was formulated after cohort B failed, so on the two cohorts above it is an explanation written with



**Figure 6:** The running covariance estimate and the streaming class, one participant per cohort (cohort A sub-01, cohort B SL01). Top row: class covariances (crosses) and the estimate's path in recording order, projected to two dimensions by multidimensional scaling on the affine-invariant Riemannian distance and coloured by the streaming class. Bottom row: distance from the estimate to each class covariance against window index, with the background shaded by the streaming class. (a,b) Cohort A,  $\tau_{\text{blk}} = 18$ ,  $m = 0.2$ : the estimate stays closer to the stop covariance throughout and approaches neither class. (c,d) Cohort B,  $\tau_{\text{blk}} = 214$ ,  $m = 0.2$ : the estimate starts near the stop covariance, moves to the walk covariance when the walk block begins, and the two distances swap back when stop returns. (e,f) Cohort B at  $m = 0.01$ : the estimate moves toward the walk covariance more slowly and crosses over once, late in the recording.



**Figure 7:** The adaptation rate condition. (a) Cohort B, five arms, three momenta, single seed; the shaded region is where the estimator memory is shorter than a class block. Measurement points are marked because the sweep is sparse. (b) The same change applied to both cohorts at matched estimator version, giving opposite signs. (c) Effect of alignment on balanced accuracy ( $align + gate$  minus  $gate$  only) against estimator memory, three data-split seeds under the current estimator, for all three cohorts; dotted lines mark each cohort's  $\tau_{blk}$ . The effect on cohort B recovers gradually as the memory passes its block length and does not become positive. No upper bound is drawn, because  $\tau_{drift}$  was not measured (Section 8). The arms shown are the components of Section 4.3 and none is the reported model; the condition is drawn here on the alignment layer because that is where it was derived, and Section 6.9 shows it governing the branch's reference, which is the estimator the reported model actually uses.

**Table 15**

Predictions registered before any model was trained on cohort C. Balanced accuracy averaged over all 20 participants, seed 0. Differences are paired within participant and given as mean  $\pm$  SD across participants.

| Prediction                | Arms compared | Difference         | Threshold | Outcome |
|---------------------------|---------------|--------------------|-----------|---------|
| P1 alignment costs little | 0.783 / 0.785 | +0.002 $\pm$ 0.039 | > -0.05   | holds   |
| P2 slowing does not help  | 0.785 / 0.779 | -0.006 $\pm$ 0.033 | < +0.05   | holds   |

their results in view. To test it in advance, predictions for a third cohort were committed to version control before any model was trained on it. The cohort, participants, data split, preprocessing, compared arms and decision thresholds were all fixed at that point.

Cohort C alternates rest and movement about every 4 s. From the published protocol alone a class block holds about five windows, and the loader later measured a median of 5.0. Against an estimator memory of 160 windows at  $m = 0.2$ , the lower bound of (5) is satisfied by a factor of about thirty, so the estimate cannot converge to either class. Two predictions follow. P1: enabling alignment costs less than 0.05, where the same comparison cost 0.211 on cohort B. P2: slowing adaptation to  $m = 0.01$  gains less than 0.05, where it recovered 0.090 on cohort B. The arms compared in P1 differ only in alignment, and those in P2 only in momentum.

Both predictions hold (Table 15). Enabling alignment changes balanced accuracy by +0.002 (higher in 11 participants, lower in 9; Wilcoxon  $p = 0.68$ ), and slowing it changes accuracy by -0.006 (higher in 9, lower in 10;  $p = 0.43$ ). The intervention that recovered 0.090 on cohort B

brings no gain here, as the protocol predicted, and enabling alignment shows no harm.

Two further seeds were run after the registered test and are reported separately from it. They agree with both verdicts and sharpen one of them. Enabling alignment changes accuracy by +0.014 (higher in 13 participants, lower in 7;  $p = 0.17$ ), so it again costs nothing, and its small positive trend is not significant. Slowing adaptation costs 0.016 (lower in 15 of 20;  $p = 0.003$ ) and is negative on all three seeds, so on this cohort it carries a small cost. Read after the fact, the sign of slowing is consistent with the lower bound on every cohort. It helps where the estimate tracks the class (cohort B, +0.118 from a memory of 160 to 3200 windows over three seeds) and costs accuracy on both cohorts with short blocks (cohort A, -0.050; cohort C, -0.016). That pattern was observed on the replication seeds and formed no part of the registered prediction.

The test covers the lower bound of Condition 1 and its contrast with cohort B. It does not test the upper bound.

## 6.8. A second-order branch, and where the rate condition sends it

Table 16 reports the branch of Section 4.2 against the stem alone, three seeds everywhere, paired across participants.

On cohort C, which no design decision touched, the branch adds 0.064 and is higher in every one of the 20 participants. It reaches 0.843 against 0.783 for the strongest of the eight published decoders in the same pipeline, and lowers expected calibration error from 0.075 to 0.062. On

**Table 16**

The tangent-space branch against the convolutional stem, three seeds, paired Wilcoxon across participants. Cohort C was held out from every design decision. Cohort B is shown at the default adaptation rate and at the rate Condition 1 prescribes for its 214-window class blocks.

| Cohort                 | Stem  | Stem + branch     | $\Delta$ | Higher |
|------------------------|-------|-------------------|----------|--------|
| A ( $n=7$ )            | 0.864 | $0.909 \pm 0.004$ | +0.045   | 6/7    |
| C ( $n=20$ ), held out | 0.779 | $0.843 \pm 0.006$ | +0.064   | 20/20  |
| B ( $n=8$ ), $m=0.2$   | 0.744 | $0.568 \pm 0.005$ | -0.175   | 0/8    |
| B ( $n=8$ ), $m=0.01$  | 0.744 | $0.718 \pm 0.004$ | -0.026   | 3/8    |

**Table 17**

Where the tangent reference comes from, three seeds,  $\Delta$  against the stem alone on the same cohort. The adapting reference is the only setting that helps on the two cohorts whose class blocks are short.

| Reference              | Cohort A | Cohort B | Cohort C |
|------------------------|----------|----------|----------|
| Running (adapting)     | +0.045   | -0.175   | +0.064   |
| Frozen at fitting mean | -0.039   | -0.026   | -        |
| None (no whitening)    | -0.021   | +0.011   | +0.003   |

cohort A it adds 0.045 in 6 of 7 participants. The effect there is not significant at seven participants ( $p = 0.22$ ), and the more robust observation is stability: the seed-to-seed spread of the grand mean falls from 0.032 for the stem to 0.004 with the branch.

Cohort B inverts. At the default rate the branch reaches 0.568 against 0.744, lower in 8 of 8 participants, and sits close to the 0.500 chance line. This is the failure Condition 1 predicts rather than a separate one. That cohort presents Walk in blocks of 214 windows and is 88 % Walk; at  $m = 0.2$  the estimator memory is 160 windows, shorter than a block, so the reference converges on the Walk covariance, and whitening a Walk window by the Walk covariance removes the variance that separates the classes. Lengthening the memory to 3200 windows, which the condition prescribes for a 214-window block, returns the cohort to 0.718, indistinguishable from the stem ( $p = 0.31$ ). The same rate that makes the alignment layer safe makes the reference safe, on a mechanism introduced independently of it.

Table 17 separates the two roles of the reference. Freezing it at the fitting-split mean removes the ability to track the streaming class but not the class imbalance in the fitting data itself, and recovers cohort B only partly. Removing it altogether repairs cohort B and is null on cohort C and negative on cohort A, so the adapting reference is where the gain comes from and the no-reference variant is a cohort B patch with no support elsewhere. We therefore retain the adapting reference and let Condition 1, read from the protocol, decide where the branch is admissible. Nothing in that choice is fitted on a confirmation cohort.

One boundary on the claim. The branch helps the stem, and it does not help the configuration that already contains the alignment layer: on cohort C the same branch added to the aligned arm changes accuracy by +0.002 ( $p = 0.52$ ,

higher in 14 of 20). Both read the same second-order structure, so they are substitutes rather than complements, and the configuration we report is the stem with the branch and no alignment layer.

## 6.9. Five cohorts, and one boundary that separates them

Table 18 collects the branch against the convolutional stem on every cohort, three seeds throughout, paired across participants. Cohorts D and E were added after the method was fixed, and nothing was tuned on either.


The branch helps on four cohorts and fails on one. The cohort it fails on is the only one whose class blocks outlast a covariance update, and that comparison is available before any model is trained:  $\tau_{\text{blk}}$  comes from the label sequence and  $N$  is the batch size. On cohorts D, C, E and A a block is 1, 5, 13 and 18 windows against an update of 32, so every update already spans several blocks and no momentum can make the reference follow one class. On cohort B a block is 214 windows, an update sits inside it, and the reference converges on the Walk covariance; whitening a Walk window by the Walk covariance removes the variance separating the classes, and accuracy falls to 0.568 against a 0.500 chance line.

Lengthening the estimator memory on cohort B moves it back monotonically: 0.568 at a memory of 160 windows, 0.650 at 428 and 0.718 at 3200, against 0.744 for the stem alone. The direction the comparison predicts is visible at every point measured.

Two limits belong with this claim. The boundary was identified after cohorts A, B and C were run, so only its applications to cohorts D and E were prospective. And the boundary itself is untested: every cohort where the branch helps has  $\tau_{\text{blk}}$  between 1 and 18, the one where it fails has 214, and no cohort has a block between 32 and 160 windows, which is why the block length appeared for a time to compete with  $N$  rather than with  $\mu$ . Table 19 settles it by the experiment that would test it directly.

## 6.10. Published decoders on the two new cohorts

Table 20 places the branch against the eight published decoders on the two cohorts added last. All are trained in the same pipeline with the same splits, seeds and metric code as on cohorts A, B and C, so the comparison is within-harness. The branch leads every published decoder on both: 0.741 against 0.682 for the strongest on cohort D, and 0.949 against 0.896 on cohort E. The stem alone also clears them, so the ordering does not depend on the branch; what the branch adds is the margin. Cohort D needs a caveat the margin itself hides. Five of its eight decoders are left near chance for most participants on at least one seed, which the roughly 100 fitting windows per participant explain, so the wide gaps there measure the fragility of those decoders under shared settings rather than anything about the branch. Paired over the nine participants with Holm correction across the eight comparisons, the three that trained give +0.059 over ShallowFBCSPNet (higher in 7 of



**Figure 8:** The branch on every cohort. (a) Change in balanced accuracy against the stem alone, paired across participants with seeds averaged within participant; bars are 95 % intervals from a paired bootstrap over participants, and the annotation gives how many participants improved. Cohorts are ordered by the class-block length of their protocol. (b) The same five effects against that block length, with the estimator memory drawn as a dashed line. The branch helps on every cohort whose blocks are shorter than one update and fails on the only one whose blocks are longer. Both axes of (b) are fixed before training: the block length is read from the labels and the memory is fixed by the optimiser settings, so the separation uses no accuracy. Table 19 moves the memory directly and crosses this boundary in both directions on one cohort.

**Table 18**

The tangent-space branch against the stem alone, three seeds, paired Wilcoxon across participants.  $\tau_{\text{blk}}$  is the median single-class block in windows, read from the fitting labels before training. The estimator memory is  $\mu = N/m = 160$  windows at the settings used throughout, so a cohort is at risk when a block outlasts that memory. The regime column follows from that comparison alone and uses no accuracy.

| Cohort                   | $n$ | $\tau_{\text{blk}}$ | Regime  | Stem  | Stem + branch | $\Delta$ | Higher | $p$                  |
|--------------------------|-----|---------------------|---------|-------|---------------|----------|--------|----------------------|
| C, motor execution       | 20  | 5                   | safe    | 0.779 | 0.843         | +0.064   | 20/20  | $8.8 \times 10^{-5}$ |
| D, BCI IV-2a             | 9   | 1                   | safe    | 0.701 | 0.741         | +0.040   | 8/9    | 0.020                |
| E, exoskeleton (DECODED) | 7   | 13                  | safe    | 0.928 | 0.949         | +0.021   | 6/7    | 0.047                |
| A, exoskeleton           | 7   | 18                  | safe    | 0.864 | 0.909         | +0.045   | 6/7    | 0.219                |
| B, treadmill             | 8   | 214                 | at risk | 0.744 | 0.568         | -0.175   | 0/8    | 0.008                |

**Table 19**

The estimator memory against the class-block length, varied directly.  $\mu = N/m$  is the memory in windows, set by the number  $N$  averaged into one covariance update and the momentum  $m$ ; controls are matched on  $N$ , which changes training, while  $m$  is inert for a stem that maintains no running statistic. The branch helps wherever  $\mu$  exceeds the block length and fails wherever it does not, on both cohorts and across a 400-fold range of  $\mu$ . Cohort A at  $N = 8$ ,  $m = 0.2$  is the row that settles the formulation:  $N$  has fallen below the block length while  $\mu$  has not, and the gain survives, so the governing quantity is the memory and not the update size.

| Cohort | $\tau_{\text{blk}}$ | $N$ | $m$  | $\mu$ | $\mu > \tau_{\text{blk}}$ | $\Delta$ | Higher condition                                                   |
|--------|---------------------|-----|------|-------|---------------------------|----------|--------------------------------------------------------------------|
| A      | 18                  | 8   | 1    | 8     | no                        | -0.076   | 0/7 before data is collected                                       |
| A      | 18                  | 8   | 0.2  | 40    | yes                       | +0.048   | 6/7 Parameter counts belong beside this                            |
| A      | 18                  | 32  | 0.2  | 160   | yes                       | +0.045   | 6/7 292,943 parameters against 19,505 for the stem, so the         |
| B      | 214                 | 32  | 0.2  | 160   | no                        | -0.175   | 0/8 gains in this table cost roughly fifteen times the parameters. |
| B      | 214                 | 256 | 0.2  | 1280  | yes                       | -0.046   | 3/8 Of the published decoders, ShallowFBCSPNet and EEGNet          |
| B      | 214                 | 32  | 0.01 | 3200  | yes                       | -0.026   | 3/8 are an order of magnitude smaller again, and EEG Con-          |

**Table 20**

Published decoders on the two cohorts added last, mean and standard deviation over three seeds, same pipeline as Table 14. <sup>†</sup> marks a decoder left near chance for at least five of the nine cohort D participants on at least one seed. The mark denotes a fragile regime rather than a fixed property of a decoder: running the identical criterion on an independent repeat of these cells on a second machine moves EEGNeX across the threshold, at four near-chance participants here and five there. Cohort D supplies about 100 fitting windows per participant, the fewest of the five cohorts, and the shared hyperparameters break five of the eight there. A margin over a decoder that did not train is not a comparison, so on cohort D the load-bearing rows are the two that did: ShallowFBCSPNet, FBLightConvNet and EEGNeX.

| Decoder                   | Cohort D      | Cohort E      |
|---------------------------|---------------|---------------|
| ShallowFBCSPNet           | 0.682 ± 0.013 | 0.886 ± 0.008 |
| FBLightConvNet            | 0.674 ± 0.004 | 0.891 ± 0.022 |
| EEGNeX                    | 0.656 ± 0.015 | 0.887 ± 0.018 |
| EEGConformer <sup>†</sup> | 0.599 ± 0.077 | 0.896 ± 0.010 |
| TSception <sup>†</sup>    | 0.579 ± 0.015 | 0.833 ± 0.011 |
| EEGNet <sup>†</sup>       | 0.526 ± 0.017 | 0.700 ± 0.054 |
| Deep4Net <sup>†</sup>     | 0.525 ± 0.007 | 0.722 ± 0.035 |
| EEGTCNet <sup>†</sup>     | 0.505 ± 0.015 | 0.695 ± 0.016 |
| Stem alone                | 0.701         | 0.928         |
| Stem + branch             | <b>0.741</b>  | <b>0.949</b>  |

9,  $p = 0.039$ ), +0.085 over EEGNeX (8 of 9,  $p = 0.039$ ) and +0.067 over FBLightConvNet (9 of 9,  $p = 0.031$ ). Those are the comparisons cohort D actually supports, and all three hold after correction.

### 6.11. The reported model against the field

Table 3 is the comparison the rest of the results build toward: the reported model, the stem it is built on, and eight published decoders, on every cohort, three seeds each. Every cell was trained in the same pipeline with the same splits, seeds, preprocessing and metric code, so the columns


are internally comparable. They are not comparable with numbers from these decoders' original papers, which used different splits and different preprocessing, and we make no claim about those.

The reported model is highest on 4 of the five cohorts. On the fifth, cohort B, it is the worst entry in the column, and that is exactly the cohort where a class block outlasts the estimator memory. What makes this a usable result rather than an unreliable one is that the stem alone leads that column: 0.744 against 0.734 for the strongest published decoder there. One of the two configurations is therefore the best entry in every column of the table, and the block-length condition selects between them from the recording protocol, before any data is collected and without consulting accuracy. Parameter counts belong beside this. The reported model has 292,943 parameters against 19,505 for the stem, so the gains in this table cost roughly fifteen times the parameters. Of the published decoders, ShallowFBCSPNet and EEGNet are an order of magnitude smaller again, and EEG Conformer is larger than ours. Nothing here is a claim about efficiency.

## 7. Discussion

### 7.1. Validating a mechanism separately from performance

Domain adaptation and invariance methods are routinely justified by a distribution-level quantity, a divergence, an alignment distance or a nuisance-decodability score, with task performance assumed to follow [65–68]. This work provides five counterexamples within a single system, three of which we produced by acting on that assumption ourselves (Fig. 9).



**Figure 9:** Five cases in which a distribution-level quantity improved (right, in blue) and the task did not follow (left, in red). The final row has no accuracy change and therefore no bar: the leakage probe rises with accuracy whether or not the decoder can access the nuisance.

1. The adaptation rate was originally selected because it maximised drift reduction (on cohort A, 45.9 % at  $m = 0.2$  against 12.1 % at  $m = 0.05$  and 0.2 % at  $m = 0.01$ ).

- That choice costs 0.118 accuracy on cohort B against a memory of 3200 windows, over three seeds (Table 6).
2. The cohort in which more drift was removed (62.5 % against 45.9 %, Table 9) is the cohort in which the layer costs 0.149 accuracy over three seeds (Table 6).
  3. A verified correction to a genuine covariance-estimation bias moved decoding by +0.005 and -0.016 (Section 6.5).
  4. A prediction registered from our own diagnosis was falsified (Table 5), because the diagnosis accounted only for the harm the estimate can do.
  5. The nuisance-leakage probe rises with accuracy whether or not the decoder can access the nuisance, measuring  $R^2 = +0.863$  for one that cannot.

Item 4 generalises the rest. Each of these quantities is one-sided by construction: it scores the failure it was designed to detect and carries no term for the capability the intervention removes. Drift reduction measures how much non-stationarity a layer removes and has no term for what the layer destroys. A diagnosis inherited from such a metric inherits its blind spot, which is why our own threshold rule was correct about one direction and silent about the other until we tested it where it predicted no effect. The practical implication is that reporting a reduced divergence, increased invariance or lowered nuisance decodability is evidence about the representation alone, saying nothing about whether the method helps, even when the reduction is large and reproducible.

## 7.2. Relation to temporally correlated test-time adaptation

The vision literature has identified that test-time adaptation degrades when the test stream is temporally correlated [20–22], and proposes reservoirs and stabilised objectives as remedies. Our result is compatible with that line and adds two things specific to the blocked-protocol case. First, the relevant quantity is a ratio of timescales, measured on the stream, and which side of the threshold a protocol falls on can be predicted in advance, as the third cohort shows (Section 6.7). Second, slowing adaptation is not free. On both cohorts with short blocks it costs accuracy, so remedies that slow adaptation can exchange class tracking for a different loss, which on these cohorts could not be attributed to drift tracking directly.

## 8. Limitations

Several limitations constrain the interpretation of these results.

**Seeds and effect size.** Every cohort in Table 3 is three data-split seeds. The effect of the branch is nevertheless modest in absolute terms, +0.021 to +0.064 where it helps, and on cohort A the paired bootstrap interval spans zero ( $[-0.020, +0.109]$ ,  $p = 0.22$  at seven participants). The cohorts that carry the claim are C, with twenty participants and all twenty improving, and D and E, where the intervals exclude zero. A reader who wants a single number should

take cohort C's +0.064 and treat cohort A as consistent with it rather than as independent support.

**No paired significance tests on accuracy.** Per-participant accuracies were not retained for the single-seed component ablation and the single-seed rate tables, so those comparisons carry dispersion without paired tests. Paired tests are reported wherever per-participant rows exist: the drift results (Table 9), the temporal and normalisation controls, the artefact control (Section 6.3), the three-seed alignment comparison and the two-seed contrasts against the stem on cohort A (Section 6.2, Table 4), the comparison with published decoders on cohort C (Section 6.11) and the cohort C predictions (Section 6.7).

**Cohort size.** Seven, eight, twenty, nine and seven participants, 51 in total. This is typical for exoskeleton BCI studies [1] but small, and the between-subject SD of 0.06 to 0.13 (Table 11) is large relative to the differences between arms.

**Best-epoch evaluation.** Model selection uses the best epoch by held-out balanced accuracy, which is conventional but upward-biased on small folds. The bias applies equally to all arms and baselines within a pipeline, so comparisons within a pipeline are unaffected, but absolute values should be read as optimistic.

**Baseline training.** The published decoders were trained with the proposed decoder's settings and no per-model tuning. EEGNet, Deep4Net and EEG-TCNet stayed near chance for several cohort C participants under those settings, so their cohort C results record a training failure and are not evidence about those architectures.

**The upper bound is not measured.** Two label-free estimates of  $\tau_{\text{drift}}$  were attempted. Block-wise covariance distances within a session are dominated by class alternation where blocks are short, and on cohort A they return about 18 windows, the class-block length. Distances between session mean covariances are already at their plateau at a gap of one session for every cohort A participant, and fall with the gap for two, so a fitted timescale lies below the sampling interval and is not a measurement; cohort B's three trials give two gaps, too few for a fit. The data therefore bound the drift timescale only from above, at about one session, which does not constrain the adaptation rate. A protocol with more closely spaced recordings is needed to test the upper bound.

**Condition 1 is hypothesis-generating.** It was formulated after observing cohort B fail. One prediction from it has since been tested out-of-sample, and that test falsified the one-sided form (Table 5), which is how the two-sided form arose. The corrected form has since been tested prospectively on cohort C (Section 6.7), where both registered predictions held. That test covers the lower bound and the contrast with cohort B; the upper bound has not been tested in advance.

**Artefact.** Retraining after ICA cleaning lowers accuracy by 0.051, concentrated in three participants (Section 6.3). The losses do not follow the amount of muscle activity removed. With seven participants, and with re-referencing part of the cleaning, the control still falls short of establishing

that the accuracy is neural. ICLabel is also not trained on gait-related artefact specifically.

**Window length differs between cohorts.** Cohort B uses 2 s windows against cohort A's 4 s, so the context pathway spans a shorter duration there. This is a confound for the context ablation specifically, though not for the rate condition, which is measured within each cohort.

**Parameter cost.** The branch is 238,028 parameters and the fusion it requires a further 32,896, against 19,505 in the stem path. The reported model is therefore about fifteen times the size of the stem it improves on, for 0.02 to 0.06 of balanced accuracy. Two of the published decoders it beats are an order of magnitude smaller again. Whether the tangent vector can be projected more cheaply, for instance by reducing its dimension before the linear layer, is not tested here and is the obvious first place to look.

**The block-length boundary was found after the fact.** The comparison between  $\tau_{\text{blk}}$  and the estimator memory separates all five cohorts with nothing fitted, but it was identified after cohorts A, B and C had been run. Only its applications to cohorts D and E were prospective, and of those only cohort C's original rate prediction was registered in the repository before the run. The boundary is supported, not established, and the test that would establish it is a cohort chosen in advance because its protocol sits near the boundary.

**The boundary is untested at the boundary.** Across cohorts the boundary is straddled but not located: every cohort where the branch helps has a block of 1 to 18 windows and the one where it fails has 214, against an update of 32. Nothing is measured between 18 and 214, so across cohorts the location of the change is inferred from five points that straddle it. Within cohorts it is measured directly: varying  $\mu$  at fixed block length crosses the boundary in both directions (Table 19).

**No calibrated threshold.** Two manipulations move the effect in the predicted direction and neither restores parity. Setting the adaptation rate from the block length with a fixed margin of two was pre-registered and falsified. Enlarging the update to 256 windows against a 214-window block recovers 74 % of the loss and still costs 0.046. The condition therefore gives an ordering and a direction, not a rate a practitioner could set, and the upper side of the two-sided condition remains unmeasured on both cohorts where it was attempted.

**Cohort E tests task, not drift.** Only one participant recorded on two dates, so its split is temporal across runs within a session. It is the closest match in the paper to the target application and the weakest test of longitudinal shift, and those two facts are in tension.

**Reproducibility across machines.** The cohort D baseline sweep was run independently on two machines. Sixteen of 24 cells differ, by up to 0.022 on a single cell, while the three-seed arm means agree to 0.003 and the strongest published decoder is identical on both. Single-cell numbers anywhere in this paper therefore carry cross-machine variation wider than several of the gaps between neighbouring rows, and only seed-averaged quantities should be compared.

## 9. Conclusion

A log-power stem reads the diagonal of a window's spatial covariance and discards the rest of it. Supplying that structure through a tangent-space branch raises balanced accuracy on four of five cohorts and 51 participants, and beats eight published decoders trained in the same pipeline. It fails on the fifth, and the protocol says in advance which case a cohort is: the branch's reference is an average over whatever is currently streaming, so it survives only while its memory outlasts a class block.

That condition is the part worth building on, and it is also the part least finished. It separates five cohorts with nothing fitted, it predicted the sign on two cohorts prospectively, and enlarging the update recovers most of the loss where it is violated. It still yields no rate a practitioner could set, its upper side has never been measured, and no cohort has been run near the boundary itself. The experiment that would settle it is small and specific: choose a protocol whose class blocks fall between 18 and 214 windows, register the prediction, and run it once.

The wider point is that a label-free adaptive statistic inherits the structure of the protocol it adapts on. Test-time adaptation methods are usually evaluated on streams that are shuffled or assumed i.i.d.; EEG protocols are blocked by design, and the block length is known before a single recording is made. Reporting it alongside accuracy would let a reader see, without running anything, whether an adaptive component is operating inside its domain of validity.

## CRedit authorship contribution statement

**Esm E Moula Chowdhury Abha:** Software, Methodology, Validation, Formal analysis, Investigation, Data curation, Writing – original draft, Visualization. **Md Wahiduzzaman Suva:** Software, Methodology, Validation, Formal analysis, Investigation, Data curation, Writing – original draft, Visualization. **Muhammad Hasibur Rashid Chayon:** Supervision, Project administration, Writing – review and editing.

## Declaration of competing interest

The authors declare no known competing financial interests or personal relationships that could have appeared to influence the work reported in this paper.

## Data availability

All five datasets are publicly available. All code, configurations and result files, including those of experiments that failed, are in the accompanying repository.

## References

- [1] Y. He, D. Eguren, J. M. Azorín, R. G. Grossman, T. P. Luu, J. L. Contreras-Vidal, Brain-machine interfaces for controlling lower-limb powered robotic systems, *Journal of Neural Engineering*

- 15 (2018) 021004. URL: <https://doi.org/10.1088/1741-2552/aaa8c0>. doi:10.1088/1741-2552/aaa8c0.
- [2] E. López-Larraz, F. Trincado-Alonso, V. Rajasekaran, S. Peres-Nombela, A. J. del Ama, J. Aranda, J. Minguez, A. Gil-Agudo, L. Montesano, Control of an ambulatory exoskeleton with a brain-machine interface for spinal cord injury gait rehabilitation, *Frontiers in Neuroscience* 10 (2016) 359. URL: <https://doi.org/10.3389/fnins.2016.00359>. doi:10.3389/fnins.2016.00359.
- [3] A. H. Do, P. T. Wang, C. E. King, S. N. Chun, Z. Nenadic, Brain-computer interface controlled robotic gait orthosis, *Journal of NeuroEngineering and Rehabilitation* 10 (2013) 111. URL: <https://doi.org/10.1186/1743-0003-10-111>. doi:10.1186/1743-0003-10-111.
- [4] J. L. Contreras-Vidal, R. G. Grossman, NeuroRex: a clinical neural interface roadmap for EEG-based brain machine interfaces to a lower body robotic exoskeleton, *Annual International Conference of the IEEE Engineering in Medicine and Biology Society (EMBC)* (2013) 1579–1582. URL: <https://doi.org/10.1109/embc.2013.6609816>. doi:10.1109/embc.2013.6609816.
- [5] A. Presacco, R. Goodman, L. Forrester, J. L. Contreras-Vidal, Neural decoding of treadmill walking from noninvasive electroencephalographic signals, *Journal of Neurophysiology* 106 (2011) 1875–1887. URL: <https://doi.org/10.1152/jn.00104.2011>. doi:10.1152/jn.00104.2011.
- [6] P. Shenoy, M. Krauledat, B. Blankertz, R. P. N. Rao, K.-R. Müller, Towards adaptive classification for BCI, *Journal of Neural Engineering* 3 (2006) R13–R23. URL: <https://doi.org/10.1088/1741-2560/3/1/r02>. doi:10.1088/1741-2560/3/1/r02.
- [7] M. Krauledat, M. Tangermann, B. Blankertz, K.-R. Müller, Towards zero training for brain-computer interfacing, *PLoS ONE* 3 (2008) e2967. URL: <https://doi.org/10.1371/journal.pone.0002967>. doi:10.1371/journal.pone.0002967.
- [8] B. Blankertz, R. Tomioka, S. Lemm, M. Kawanabe, K.-R. Müller, Optimizing spatial filters for robust EEG single-trial analysis, *IEEE Signal Processing Magazine* 25 (2008) 41–56. URL: <https://doi.org/10.1109/msp.2008.4408441>. doi:10.1109/msp.2008.4408441.
- [9] C. Vidaurre, C. Sannelli, K.-R. Müller, B. Blankertz, Machine-learning-based coadaptive calibration for brain-computer interfaces, *Neural Computation* 23 (2011) 791–816. URL: [https://doi.org/10.1162/neco\\_a.00089](https://doi.org/10.1162/neco_a.00089). doi:10.1162/neco\_a.00089.
- [10] R. T. Schirrmeister, J. T. Springenberg, L. D. J. Fiederer, M. Glasstetter, K. Eggersperger, M. Tangermann, F. Hutter, W. Burgard, T. Ball, Deep learning with convolutional neural networks for EEG decoding and visualization, *Human Brain Mapping* 38 (2017) 5391–5420. URL: <https://doi.org/10.1002/hbm.23730>. doi:10.1002/hbm.23730.
- [11] A. Barachant, S. Bonnet, M. Congedo, C. Jutten, Multiclass brain-computer interface classification by Riemannian geometry, *IEEE Transactions on Biomedical Engineering* 59 (2012) 920–928. URL: <https://doi.org/10.1109/tbme.2011.2172210>. doi:10.1109/tbme.2011.2172210.
- [12] M. Congedo, A. Barachant, R. Bhatia, Riemannian geometry for EEG-based brain-computer interfaces: a primer and a review, *Brain-Computer Interfaces* 4 (2017) 155–174. URL: <https://doi.org/10.1080/2326263x.2017.1297192>. doi:10.1080/2326263x.2017.1297192.
- [13] H. He, D. Wu, Transfer learning for brain-computer interfaces: a Euclidean space data alignment approach, *IEEE Transactions on Biomedical Engineering* 67 (2020) 399–410. URL: <https://doi.org/10.1109/tbme.2019.2913914>. doi:10.1109/tbme.2019.2913914.
- [14] P. L. C. Rodrigues, C. Jutten, M. Congedo, Riemannian Procrustes analysis: transfer learning for brain-computer interfaces, *IEEE Transactions on Biomedical Engineering* 66 (2019) 2390–2401. URL: <https://doi.org/10.1109/tbme.2018.2889705>. doi:10.1109/tbme.2018.2889705.
- [15] O. Yair, M. Ben-Chen, R. Talmon, Parallel transport on the cone manifold of SPD matrices for domain adaptation, *IEEE Transactions on Signal Processing* 67 (2019) 1797–1811. URL: <https://doi.org/10.1109/tsp.2019.2894801>. doi:10.1109/tsp.2019.2894801.
- [16] P. Zanini, M. Congedo, C. Jutten, S. Said, Y. Berthoumieu, Transfer learning: a Riemannian geometry framework with applications to brain-computer interfaces, *IEEE Transactions on Biomedical Engineering* 65 (2018) 1107–1116. URL: <https://doi.org/10.1109/tbme.2017.2742541>. doi:10.1109/tbme.2017.2742541.
- [17] D. Wang, E. Shelhamer, S. Liu, B. Olshausen, T. Darrell, Tent: fully test-time adaptation by entropy minimization, in: *International Conference on Learning Representations (ICLR)*, 2021. URL: <https://arxiv.org/abs/2006.10726>.
- [18] S. Schneider, E. Rusak, L. Eck, O. Bringmann, W. Brendel, M. Bethge, Improving robustness against common corruptions by covariate shift adaptation, in: *Advances in Neural Information Processing Systems (NeurIPS)*, 2020. URL: <https://arxiv.org/abs/2006.16971>.
- [19] Z. Nado, S. Padhy, D. Sculley, A. D'Amour, B. Lakshminarayanan, J. Snoek, Evaluating prediction-time batch normalization for robustness under covariate shift, *arXiv preprint arXiv:2006.10963* (2020). URL: <https://arxiv.org/abs/2006.10963>.
- [20] T. Gong, J. Jeong, T. Kim, Y. Kim, J. Shin, S.-J. Lee, NOTE: robust continual test-time adaptation against temporal correlation, in: *Advances in Neural Information Processing Systems (NeurIPS)*, 2022. URL: <https://doi.org/10.52202/068431-1976>. doi:10.52202/068431-1976.
- [21] S. Niu, J. Wu, Y. Zhang, Z. Wen, Y. Chen, P. Zhao, M. Tan, Towards stable test-time adaptation in dynamic wild world, in: *International Conference on Learning Representations (ICLR)*, 2023. URL: <https://arxiv.org/abs/2302.12400>.
- [22] Q. Wang, O. Fink, L. Van Gool, D. Dai, Continual test-time domain adaptation, in: *IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2022. URL: <https://doi.org/10.1109/cvpr52688.2022.00706>. doi:10.1109/cvpr52688.2022.00706.
- [23] V. J. Lawhern, A. J. Solon, N. R. Waytowich, S. M. Gordon, C. P. Hung, B. J. Lance, EEGNet: a compact convolutional neural network for EEG-based brain-computer interfaces, *Journal of Neural Engineering* 15 (2018) 056013. URL: <https://doi.org/10.1088/1741-2552/aace8c>. doi:10.1088/1741-2552/aace8c.
- [24] X. Chen, X. Teng, H. Chen, Y. Pan, P. Geyer, Toward reliable signals decoding for electroencephalogram: a benchmark study to EEGNeX, *Biomedical Signal Processing and Control* 87 (2024) 105475. URL: <https://doi.org/10.1016/j.bspc.2023.105475>. doi:10.1016/j.bspc.2023.105475.
- [25] Y. Roy, H. Banville, I. Albuquerque, A. Gramfort, T. H. Falk, J. Faubert, Deep learning-based electroencephalography analysis: a systematic review, *Journal of Neural Engineering* 16 (2019) 051001. URL: <https://doi.org/10.1088/1741-2552/ab260c>. doi:10.1088/1741-2552/ab260c.
- [26] A. Craik, Y. He, J. L. Contreras-Vidal, Deep learning for electroencephalogram (EEG) classification tasks: a review, *Journal of Neural Engineering* 16 (2019) 031001. URL: <https://doi.org/10.1088/1741-2552/ab0ab5>. doi:10.1088/1741-2552/ab0ab5.
- [27] Y. Song, Q. Zheng, B. Liu, X. Gao, EEG Conformer: convolutional transformer for EEG decoding and visualization, *IEEE Transactions on Neural Systems and Rehabilitation Engineering* 31 (2023) 710–719. URL: <https://doi.org/10.1109/tnsre.2022.3230250>. doi:10.1109/tnsre.2022.3230250.
- [28] T. M. Ingolfsson, M. Hersche, X. Wang, N. Kobayashi, L. Cavigelli, L. Benini, EEG-TCNet: an accurate temporal convolutional network for embedded motor-imagery brain-machine interfaces, in: *IEEE International Conference on Systems, Man, and Cybernetics (SMC)*, 2020. URL: <https://doi.org/10.1109/smc42975.2020.9283028>. doi:10.1109/smc42975.2020.9283028.
- [29] A. Barachant, S. Bonnet, M. Congedo, C. Jutten, Classification of covariance matrices using a Riemannian-based kernel for BCI applications, *Neurocomputing* 112 (2013) 172–178. URL: <https://doi.org/10.1016/j.neucom.2012.12.039>. doi:10.1016/j.neucom.2012.12.039.
- [30] F. Lotte, L. Bougrain, A. Cichocki, M. Clerc, M. Congedo, A. Rakotomamonjy, F. Yger, A review of classification algorithms for EEG-based brain-computer interfaces: a 10 year update, *Journal of Neural Engineering* 15 (2018) 031005. URL: <https://doi.org/10.1088/1741-2552/ab0ab5>.

- 1088/1741-2552/aab2f2. doi:10.1088/1741-2552/aab2f2.
- [31] D. Wu, Y. Xu, B.-L. Lu, Transfer learning for EEG-based brain-computer interfaces: a review of progress made since 2016, *IEEE Transactions on Cognitive and Developmental Systems* 14 (2022) 4–19. URL: <https://doi.org/10.1109/tcds.2020.3007453>. doi:10.1109/tcds.2020.3007453.
  - [32] Z. Huang, L. Van Gool, A Riemannian network for SPD matrix learning, in: *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 31, 2017.
  - [33] Y. Sun, X. Wang, Z. Liu, J. Miller, A. A. Efros, M. Hardt, Test-time training with self-supervision for generalization under distribution shifts, in: *Proceedings of the 37th International Conference on Machine Learning (ICML)*, 2020. URL: <https://www.semanticscholar.org/paper/a436e41141a9d9c5dfc44d16fade67b9a3b76778>.
  - [34] M. Boudiaf, R. Mueller, I. Ben Ayed, L. Bertinetto, Parameter-free on-line test-time adaptation, in: *IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2022. URL: <https://doi.org/10.1109/cvpr52688.2022.00816>. doi:10.1109/cvpr52688.2022.00816.
  - [35] M. P. Naeini, G. F. Cooper, M. Hauskrecht, Obtaining well calibrated probabilities using Bayesian binning, in: *AAAI Conference on Artificial Intelligence*, 2015. URL: <https://doi.org/10.1609/aaai.v29i1.9602>. doi:10.1609/aaai.v29i1.9602.
  - [36] C. Guo, G. Pleiss, Y. Sun, K. Q. Weinberger, On calibration of modern neural networks, in: *Proceedings of the 34th International Conference on Machine Learning (ICML)*, 2017. URL: <https://arxiv.org/abs/1706.04599>.
  - [37] Y. Ovadia, E. Fertig, J. Ren, Z. Nado, D. Sculley, S. Nowozin, J. V. Dillon, B. Lakshminarayanan, J. Snoek, Can you trust your model's uncertainty? Evaluating predictive uncertainty under dataset shift, in: *Advances in Neural Information Processing Systems (NeurIPS)*, 2019. URL: <https://arxiv.org/abs/1906.02530>.
  - [38] C. K. Chow, On optimum recognition error and reject tradeoff, *IEEE Transactions on Information Theory* 16 (1970) 41–46. URL: <https://doi.org/10.1109/tit.1970.1054406>. doi:10.1109/tit.1970.1054406.
  - [39] R. El-Yaniv, Y. Wiener, On the foundations of noise-free selective classification, *Journal of Machine Learning Research* 11 (2010) 1605–1641.
  - [40] Y. Geifman, R. El-Yaniv, SelectiveNet: a deep neural network with an integrated reject option, in: *Proceedings of the 36th International Conference on Machine Learning (ICML)*, 2019, pp. 2151–2159. URL: <https://arxiv.org/abs/1901.09192>.
  - [41] Y. Geifman, R. El-Yaniv, Selective classification for deep neural networks, in: *Advances in Neural Information Processing Systems (NIPS)*, 2017. URL: <https://arxiv.org/abs/1705.08500>.
  - [42] J. T. Gwin, K. Gramann, S. Makeig, D. P. Ferris, Electroocortical activity is coupled to gait cycle phase during treadmill walking, *NeuroImage* 54 (2011) 1289–1296. URL: <https://doi.org/10.1016/j.neuroimage.2010.08.066>. doi:10.1016/j.neuroimage.2010.08.066.
  - [43] M. Seeber, R. Scherer, J. Wagner, T. Solis-Escalante, G. R. Müller-Putz, EEG beta suppression and low gamma modulation are different elements of human upright walking, *Frontiers in Human Neuroscience* 8 (2014) 485. URL: <https://doi.org/10.3389/fnhum.2014.00485>. doi:10.3389/fnhum.2014.00485.
  - [44] T. P. Luu, S. Nakagome, Y. He, J. L. Contreras-Vidal, Real-time EEG-based brain-computer interface to a virtual avatar enhances cortical involvement in human treadmill walking, *Scientific Reports* 7 (2017) 8895. URL: <https://doi.org/10.1038/s41598-017-09187-0>. doi:10.1038/s41598-017-09187-0.
  - [45] J. Wagner, T. Solis-Escalante, P. Grieshofer, C. Neuper, G. Müller-Putz, R. Scherer, Level of participation in robotic-assisted treadmill walking modulates midline sensorimotor EEG rhythms in able-bodied subjects, *NeuroImage* 63 (2012) 1203–1211. URL: <https://doi.org/10.1016/j.neuroimage.2012.08.019>. doi:10.1016/j.neuroimage.2012.08.019.
  - [46] G. Pfurtscheller, F. H. Lopes da Silva, Event-related EEG/MEG synchronization and desynchronization: basic principles, *Clinical Neurophysiology* 110 (1999) 1842–1857. URL: [https://doi.org/10.1016/s1388-2457\(99\)00141-8](https://doi.org/10.1016/s1388-2457(99)00141-8). doi:10.1016/s1388-2457(99)00141-8.
  - [47] G. Pfurtscheller, C. Brunner, A. Schlögl, F. H. Lopes da Silva, Mu rhythm (de)synchronization and EEG single-trial classification of different motor imagery tasks, *NeuroImage* 31 (2006) 153–159. URL: <https://doi.org/10.1016/j.neuroimage.2005.12.003>. doi:10.1016/j.neuroimage.2005.12.003.
  - [48] M. Sugiyama, M. Krauledat, K.-R. Müller, Covariate shift adaptation by importance weighted cross validation, *Journal of Machine Learning Research* 8 (2007) 985–1005. URL: <https://doi.org/10.5555/1314498.1390324>. doi:10.5555/1314498.1390324.
  - [49] S. Ioffe, C. Szegedy, Batch normalization: accelerating deep network training by reducing internal covariate shift, in: *Proceedings of the 32nd International Conference on Machine Learning (ICML)*, 2015, pp. 448–456. URL: <https://arxiv.org/abs/1502.03167>.
  - [50] A. Vaswani, N. Shazeer, N. Parmar, J. Uszkoreit, L. Jones, A. N. Gomez, L. Kaiser, I. Polosukhin, Attention is all you need, in: *Advances in Neural Information Processing Systems (NIPS)*, 2017. URL: <https://arxiv.org/abs/1706.03762>.
  - [51] H. Phan, F. Andreotti, N. Cooray, O. Y. Chén, M. De Vos, SeqSleepNet: end-to-end hierarchical recurrent neural network for sequence-to-sequence automatic sleep staging, *IEEE Transactions on Neural Systems and Rehabilitation Engineering* 27 (2019) 400–410. URL: <https://doi.org/10.1109/tnsre.2019.2896659>. doi:10.1109/tnsre.2019.2896659.
  - [52] H. Phan, O. Y. Chén, M. C. Tran, P. Koch, A. Mertins, M. De Vos, XSleepNet: multi-view sequential model for automatic sleep staging, *IEEE Transactions on Pattern Analysis and Machine Intelligence* 44 (2021) 5903–5915. URL: <https://doi.org/10.1109/tpami.2021.3070057>. doi:10.1109/tpami.2021.3070057.
  - [53] A. Supratak, H. Dong, C. Wu, Y. Guo, DeepSleepNet: a model for automatic sleep stage scoring based on raw single-channel EEG, *IEEE Transactions on Neural Systems and Rehabilitation Engineering* 25 (2017) 1998–2008. URL: <https://doi.org/10.1109/tnsre.2017.2721116>. doi:10.1109/tnsre.2017.2721116.
  - [54] G. Schalk, D. McFarland, T. Hinterberger, N. Birbaumer, J. Wolpaw, BCI2000: A general-purpose brain-computer interface (BCI) system, *IEEE Transactions on Biomedical Engineering* 51 (2004) 1034–1043. URL: <https://doi.org/10.1109/tbme.2004.827072>. doi:10.1109/tbme.2004.827072.
  - [55] A. L. Goldberger, L. A. N. Amaral, L. Glass, J. M. Hausdorff, P. C. Ivanov, R. G. Mark, J. E. Mietus, G. B. Moody, C.-K. Peng, H. E. Stanley, Physiobank, PhysioToolkit, and physionet, *Circulation* 101 (2000). URL: <https://doi.org/10.1161/01.cir.101.23.e215>. doi:10.1161/01.cir.101.23.e215.
  - [56] A. Gramfort, M. Luessi, E. Larson, D. A. Engemann, D. Strohmeier, C. Brodbeck, R. Goj, M. Jas, T. Brooks, L. Parkkonen, M. Hämäläinen, MEG and EEG data analysis with MNE-Python, *Frontiers in Neuroscience* 7 (2013) 267. URL: <https://doi.org/10.3389/fnins.2013.00267>. doi:10.3389/fnins.2013.00267.
  - [57] D. P. Kingma, J. Ba, Adam: a method for stochastic optimization, in: *International Conference on Learning Representations (ICLR)*, 2015. URL: <https://arxiv.org/abs/1412.6980>.
  - [58] R. T. Schirrmeister, L. Gemein, K. Eggensperger, F. Hutter, T. Ball, Deep learning with convolutional neural networks for decoding and visualization of EEG pathology, *IEEE Signal Processing in Medicine and Biology Symposium (SPMB)* (2017). URL: <https://doi.org/10.1109/spmb.2017.8257015>. doi:10.1109/spmb.2017.8257015.
  - [59] Y. Ding, N. Robinson, S. Zhang, Q. Zeng, C. Guan, TSception: capturing temporal dynamics and spatial asymmetry from EEG for emotion recognition, *IEEE Transactions on Affective Computing* 14 (2023) 2238–2250. URL: <https://doi.org/10.1109/taffc.2022.3169001>. doi:10.1109/taffc.2022.3169001.
  - [60] A. Salami, J. Andreu-Perez, H. Gillmeister, EEG-ITNet: an explainable inception temporal convolutional network for motor imagery classification, *IEEE Access* 10 (2022) 36672–36685. URL: <https://doi.org/10.1109/access.2022.3161489>. doi:10.1109/access.2022.3161489.

- [61] P. Ablin, J.-F. Cardoso, A. Gramfort, Faster independent component analysis by preconditioning with Hessian approximations, *IEEE Transactions on Signal Processing* 66 (2018) 4040–4049. URL: <https://doi.org/10.1109/tsp.2018.2844203>. doi:10.1109/tsp.2018.2844203.
- [62] L. Pion-Tonachini, K. Kreutz-Delgado, S. Makeig, ICLabel: An automated electroencephalographic independent component classifier, dataset, and website, *NeuroImage* 198 (2019) 181–197. URL: <https://doi.org/10.1016/j.neuroimage.2019.05.026>. doi:10.1016/j.neuroimage.2019.05.026.
- [63] T. Castermans, M. Duvinage, G. Cheron, T. Dutoit, About the cortical origin of the low-delta and high-gamma rhythms observed in EEG signals during treadmill walking, *Neuroscience Letters* 561 (2014) 166–170. URL: <https://doi.org/10.1016/j.neulet.2013.12.059>. doi:10.1016/j.neulet.2013.12.059.
- [64] J. E. Kline, H. J. Huang, K. L. Snyder, D. P. Ferris, Isolating gait-related movement artifacts in electroencephalography during human walking, *Journal of Neural Engineering* 12 (2015) 046022. URL: <https://doi.org/10.1088/1741-2560/12/4/046022>. doi:10.1088/1741-2560/12/4/046022.
- [65] Y. Ganin, E. Ustinova, H. Ajakan, P. Germain, H. Larochelle, F. Laviolette, M. Marchand, V. Lempitsky, Domain-adversarial training of neural networks, *Journal of Machine Learning Research* 17 (2016) 1–35. URL: [https://doi.org/10.1007/978-3-319-58347-1\\_10](https://doi.org/10.1007/978-3-319-58347-1_10). doi:10.1007/978-3-319-58347-1\_10.
- [66] M. Long, Y. Cao, J. Wang, M. I. Jordan, Learning transferable features with deep adaptation networks, in: *Proceedings of the 32nd International Conference on Machine Learning (ICML)*, 2015. URL: <https://arxiv.org/abs/1502.02791>.
- [67] B. Sun, K. Saenko, Deep CORAL: correlation alignment for deep domain adaptation, in: *ECCV Workshops*, 2016. URL: [https://doi.org/10.1007/978-3-319-49409-8\\_35](https://doi.org/10.1007/978-3-319-49409-8_35). doi:10.1007/978-3-319-49409-8\_35.
- [68] S. Ben-David, J. Blitzer, K. Crammer, A. Kulesza, F. Pereira, J. W. Vaughan, A theory of learning from different domains, *Machine Learning* 79 (2010) 151–175. URL: <https://doi.org/10.1007/s10994-009-5152-4>. doi:10.1007/s10994-009-5152-4.